

A photograph of a modern brick building with multiple stories and numerous windows. In the foreground, there is a lawn with a variety of colorful flowers, primarily yellow. A paved walkway leads towards the building.

# Databases & SQL

---

**CSCI 3308**  
**Liz Boese**

# Objectives

---

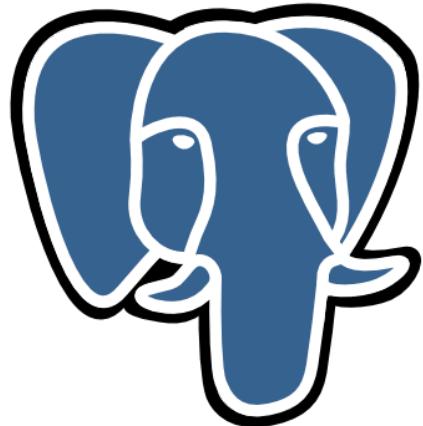
- ◆ Database Basics
- ◆ SQL
- ◆ NoSQL

---

# **RELATIONAL DATABASES**

# Databases

**ORACLE®**  
DATABASE



Microsoft®  
**SQL Server®**  
**SQLite**

PostgreSQL

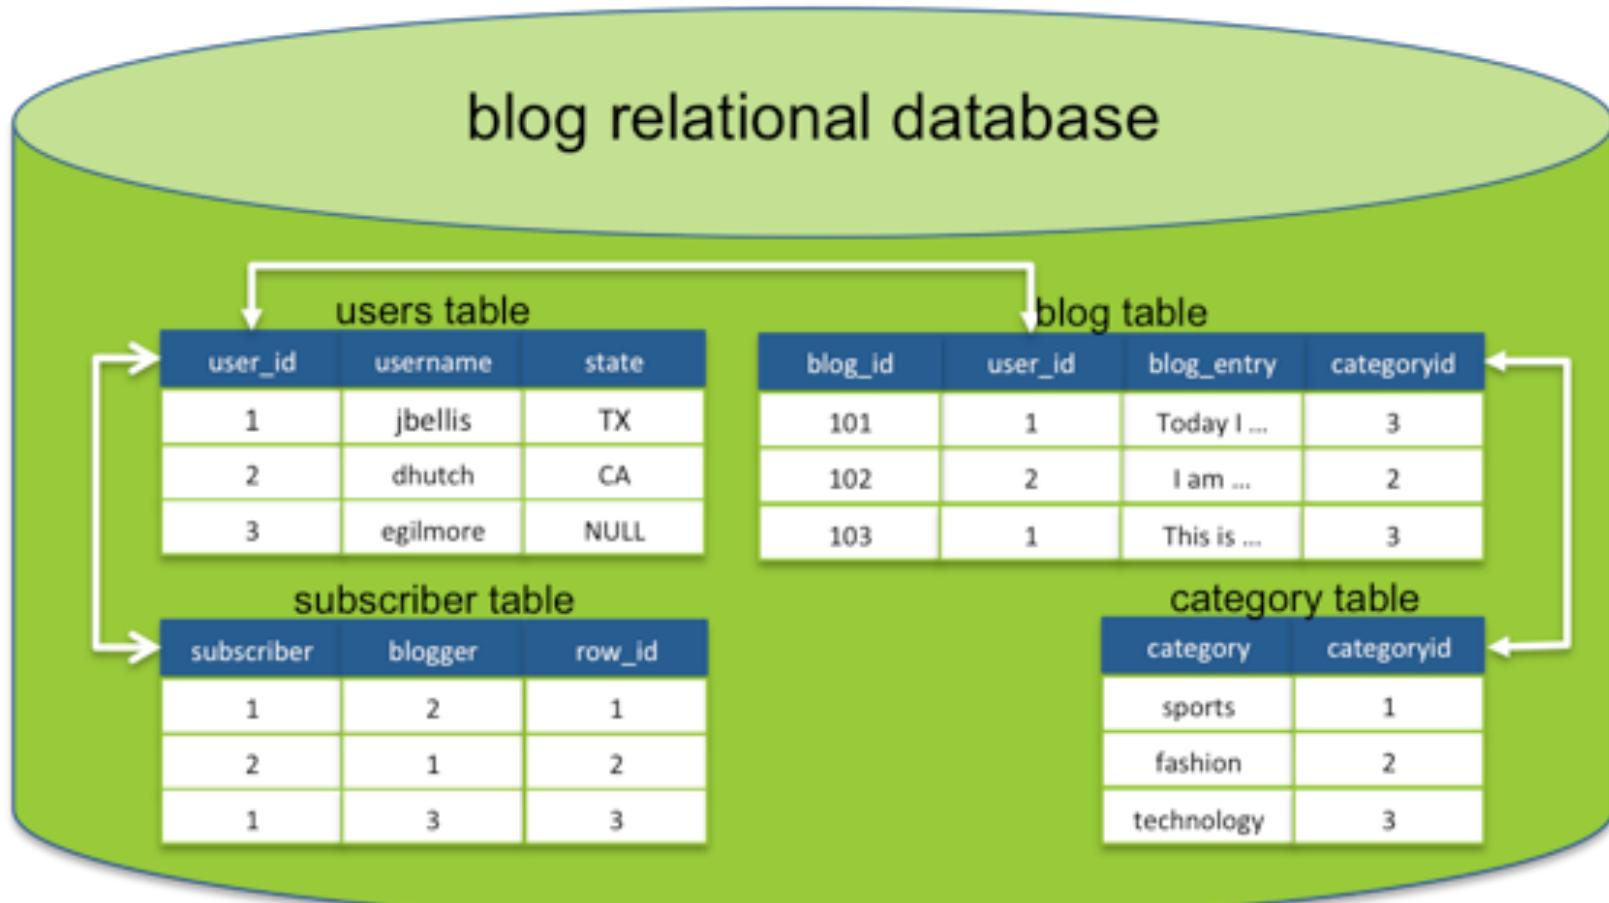
**MySQL®**

We won't be able  
to deliver our product  
in time because of  
some issue with MySQL...



# Relational Databases

- ◆ Relational Databases store data according to a set of defined entities and relations
- ◆ Entities (tables) consist of related attributes (fields)



# Entity Data Stored in Tables

---

- ◆ Table columns contain attributes
- ◆ Table rows (tuples) contain items

**Table:** Category

ID Category	Name	Description
1	Auto Parts	Things to service a car
2	Electronics	TVs, DVD players, etc.

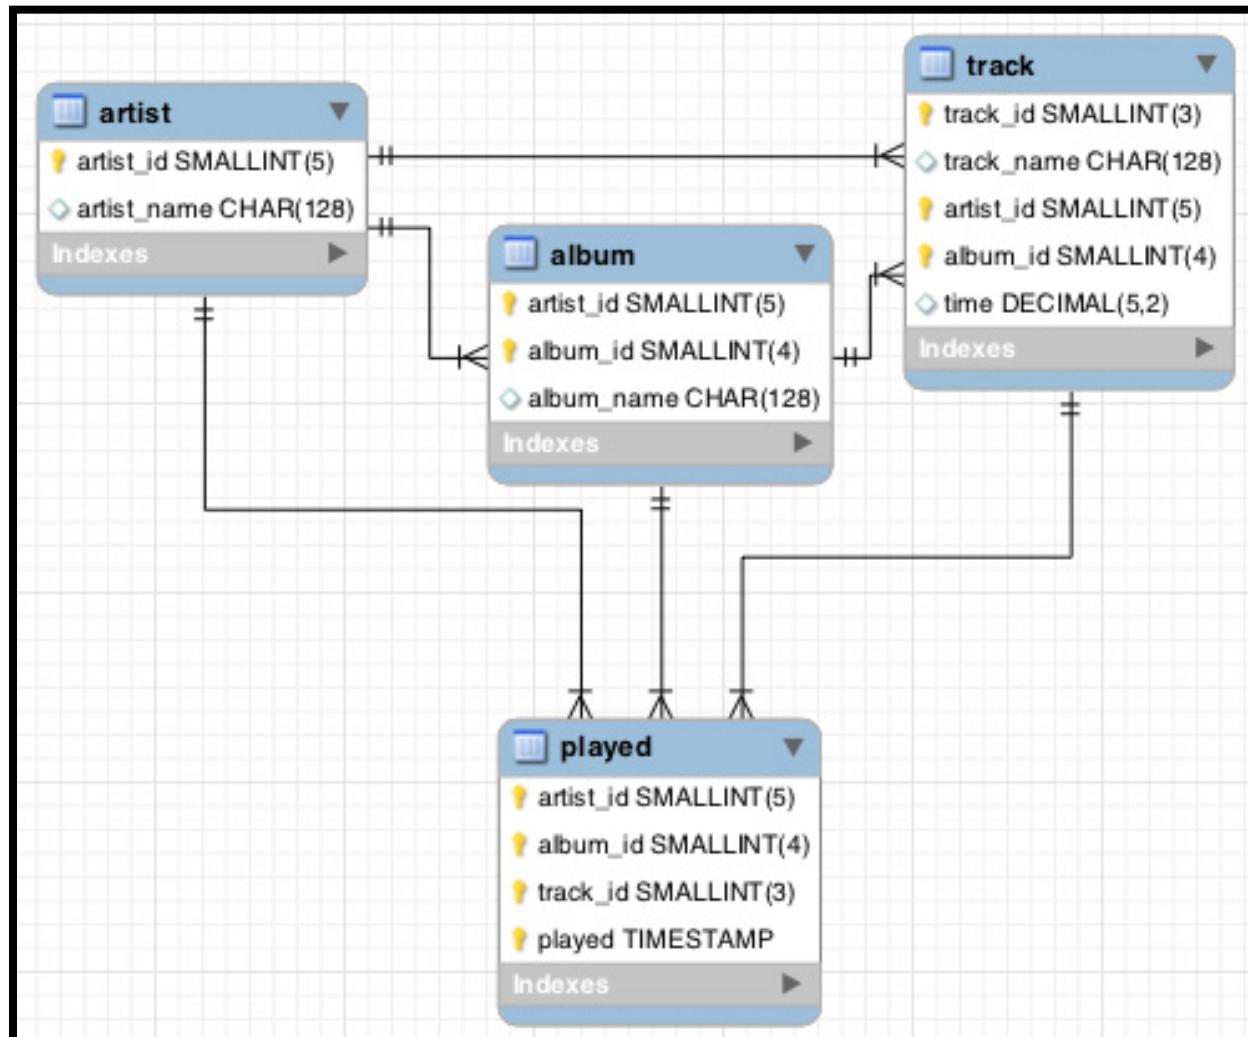
# ER Diagrams

- ◆ ER Diagrams support modeling and analysis of database relational structure

- **Relations**

- depict how two entities are related:

- “X entity contains zero or more of Y entity”*



## Definitions:

**entity** something about which data is collected, stored, and maintained

**attribute** a characteristic of an entity

**relationship** an association between entities

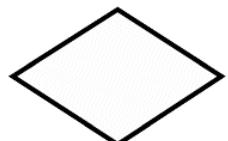
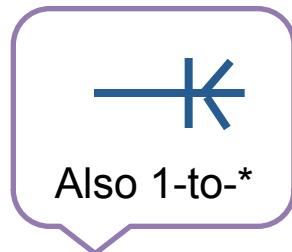
**entity type** a class of entities that have the same set of attributes

**record** an ordered set of attribute values that describe an instance of an entity type

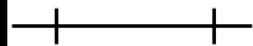
## Symbols:



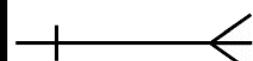
entity type



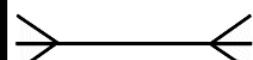
relationship between entities



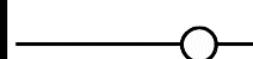
one-to-one association



one-to-many association



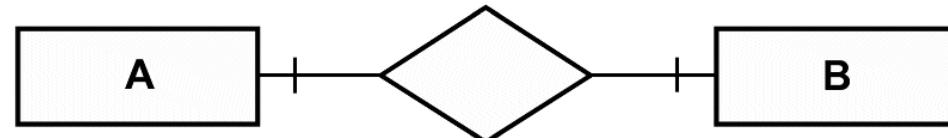
many-to-many association



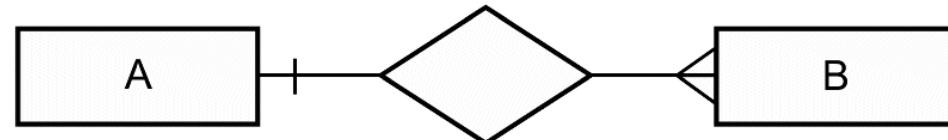
partly optional association

## Examples:

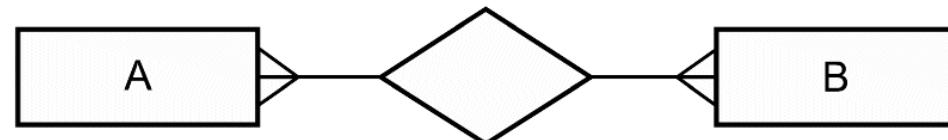
One A is associated with one B:



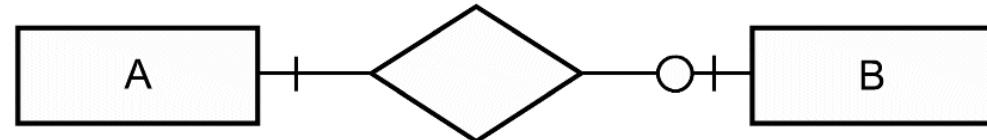
One A is associated with one or more B's:



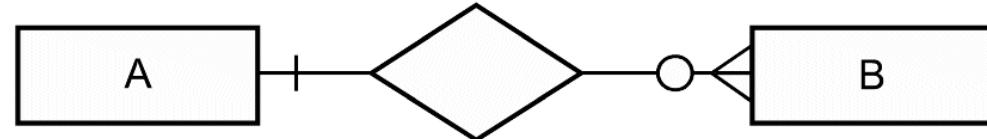
One or more A's are associated with one or more B's:



One A is associated with zero or one B:



One A is associated with zero or more B's:

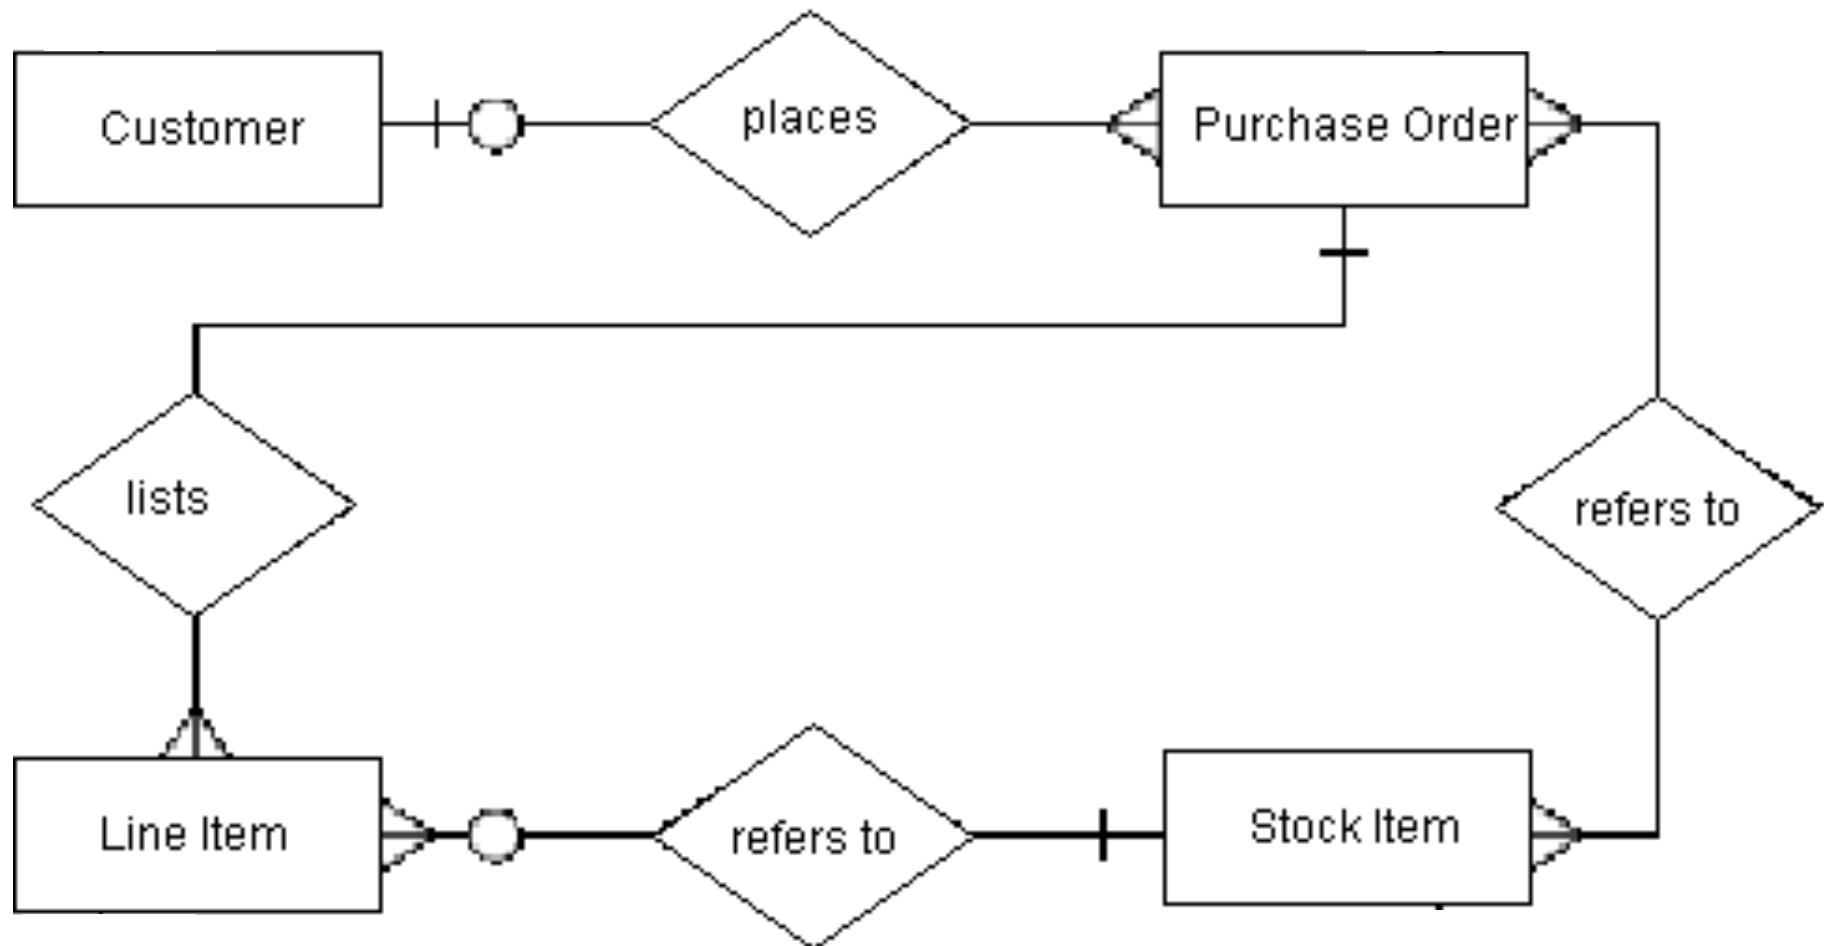


One A is associated with one B and one C:

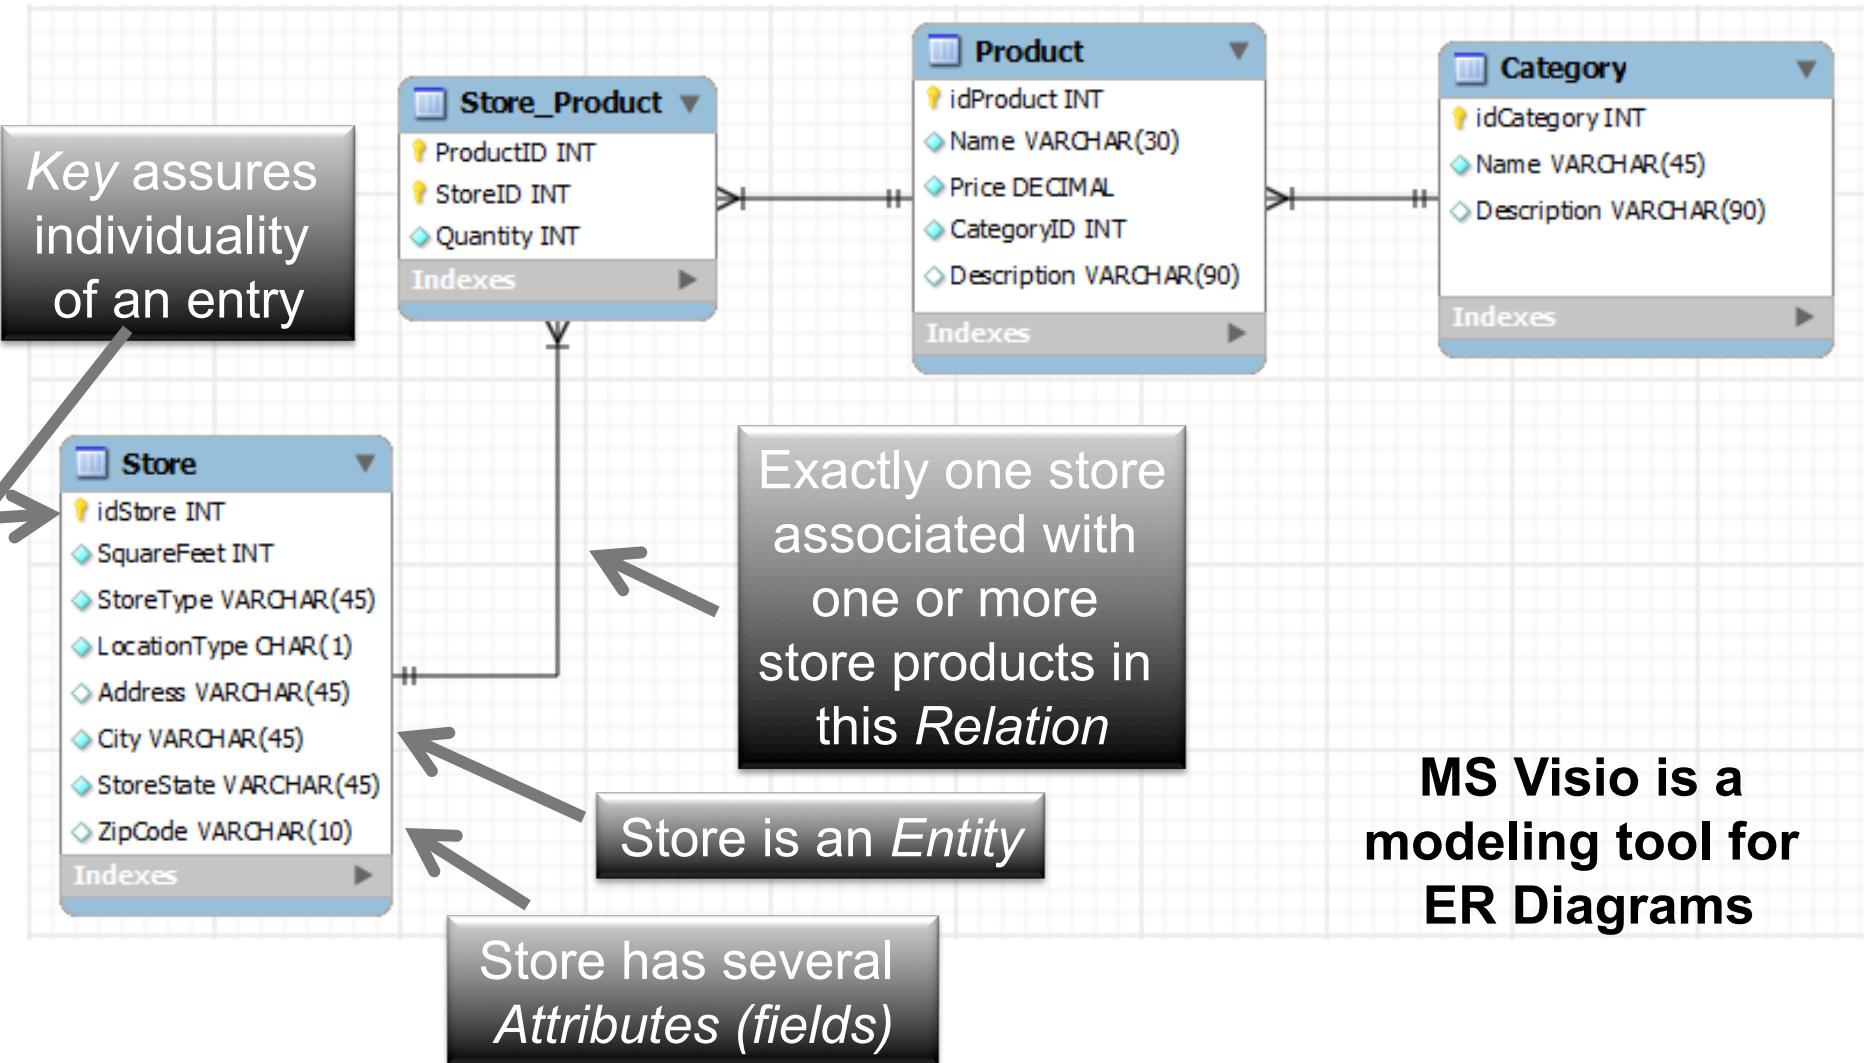


# ER Diagram

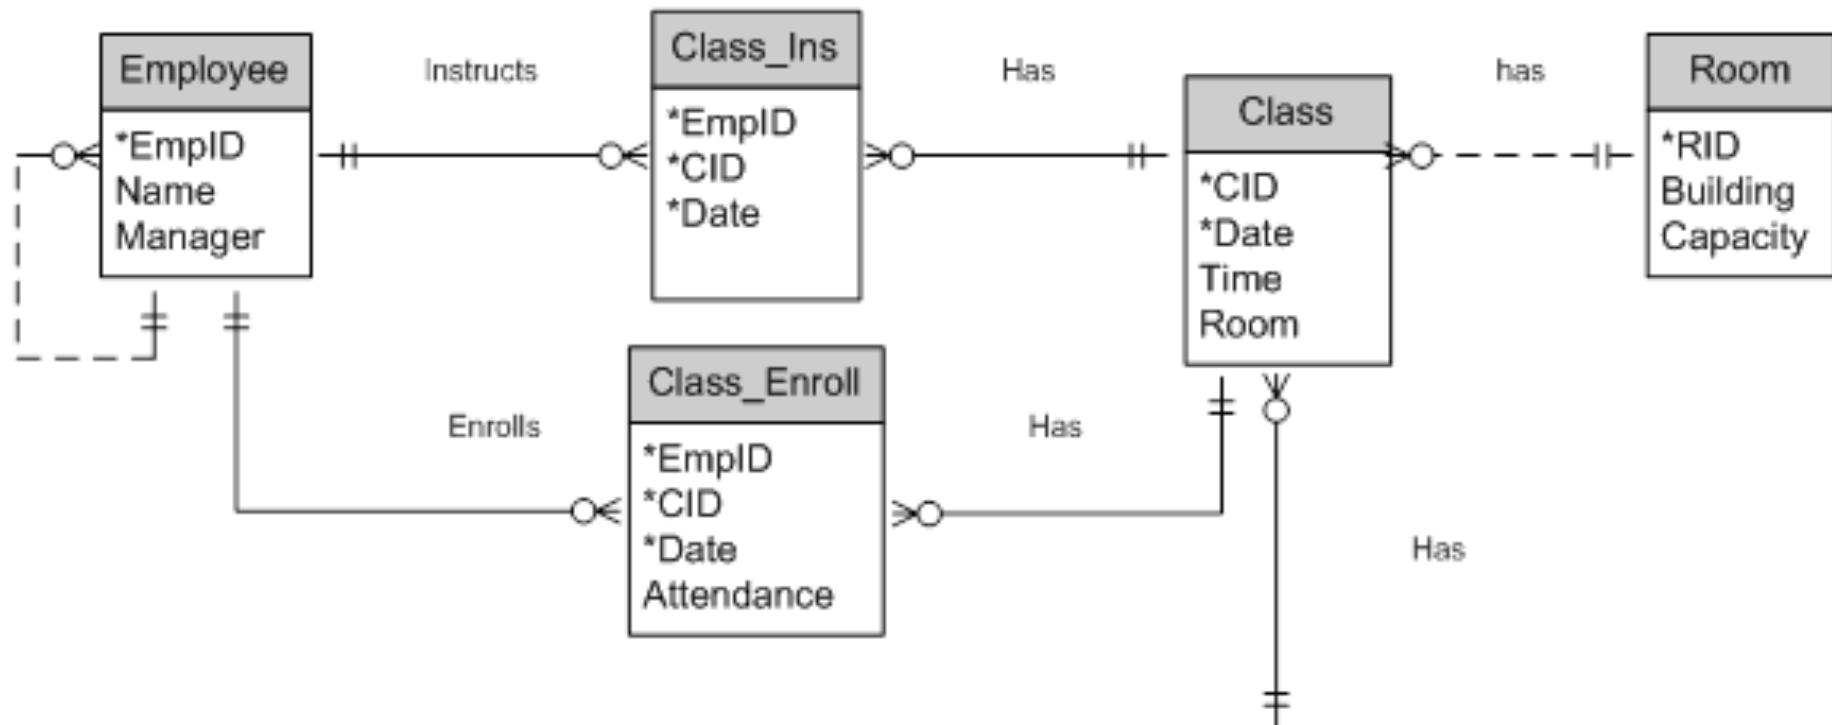
---



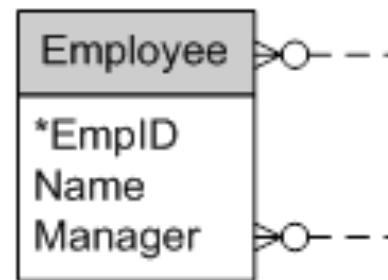
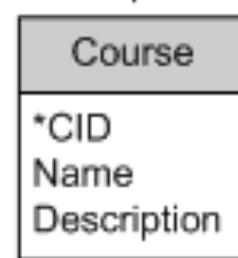
# ER Diagram

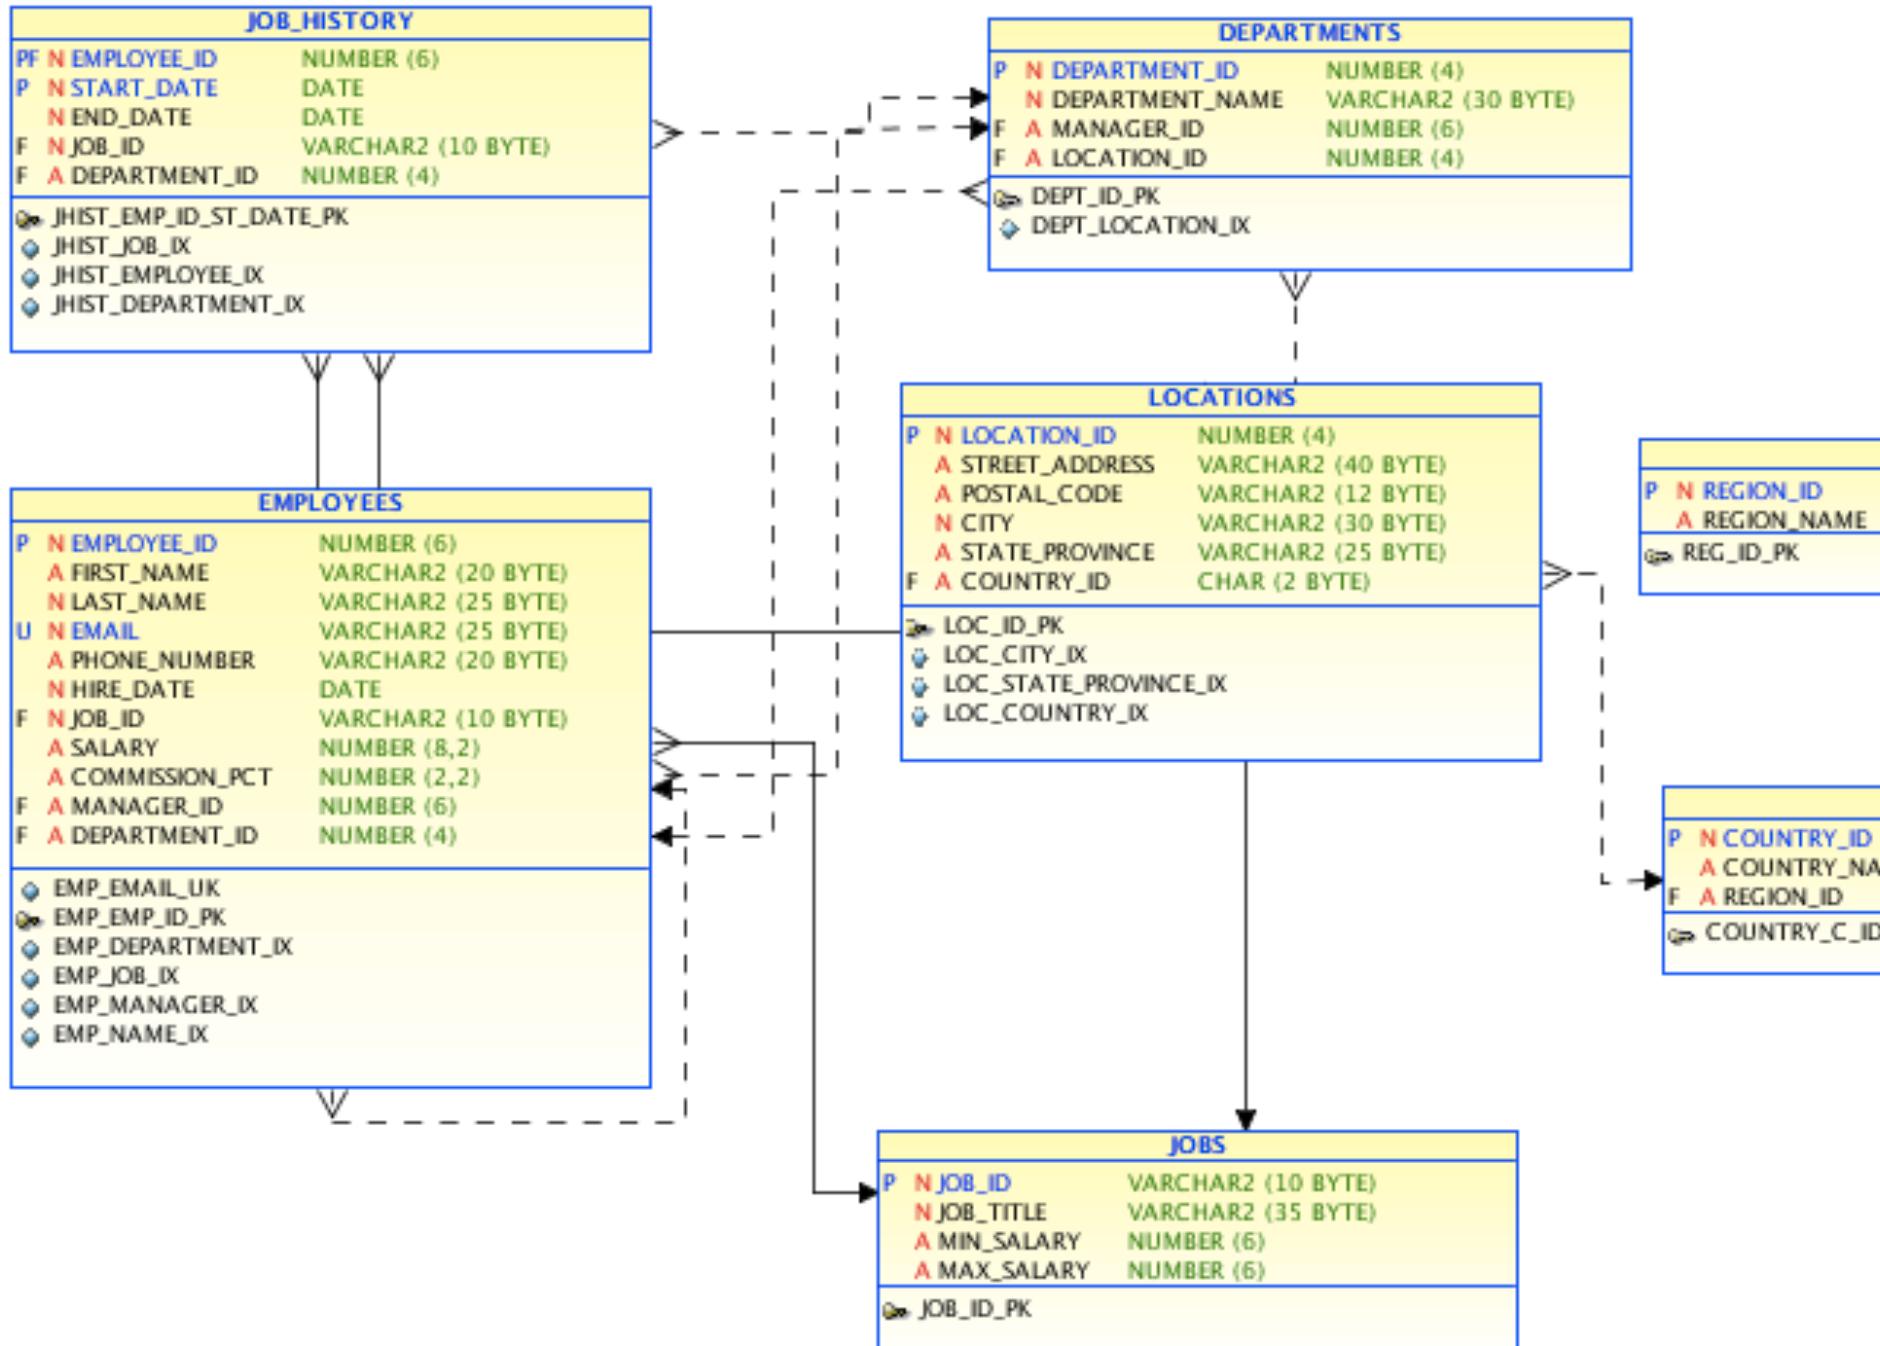


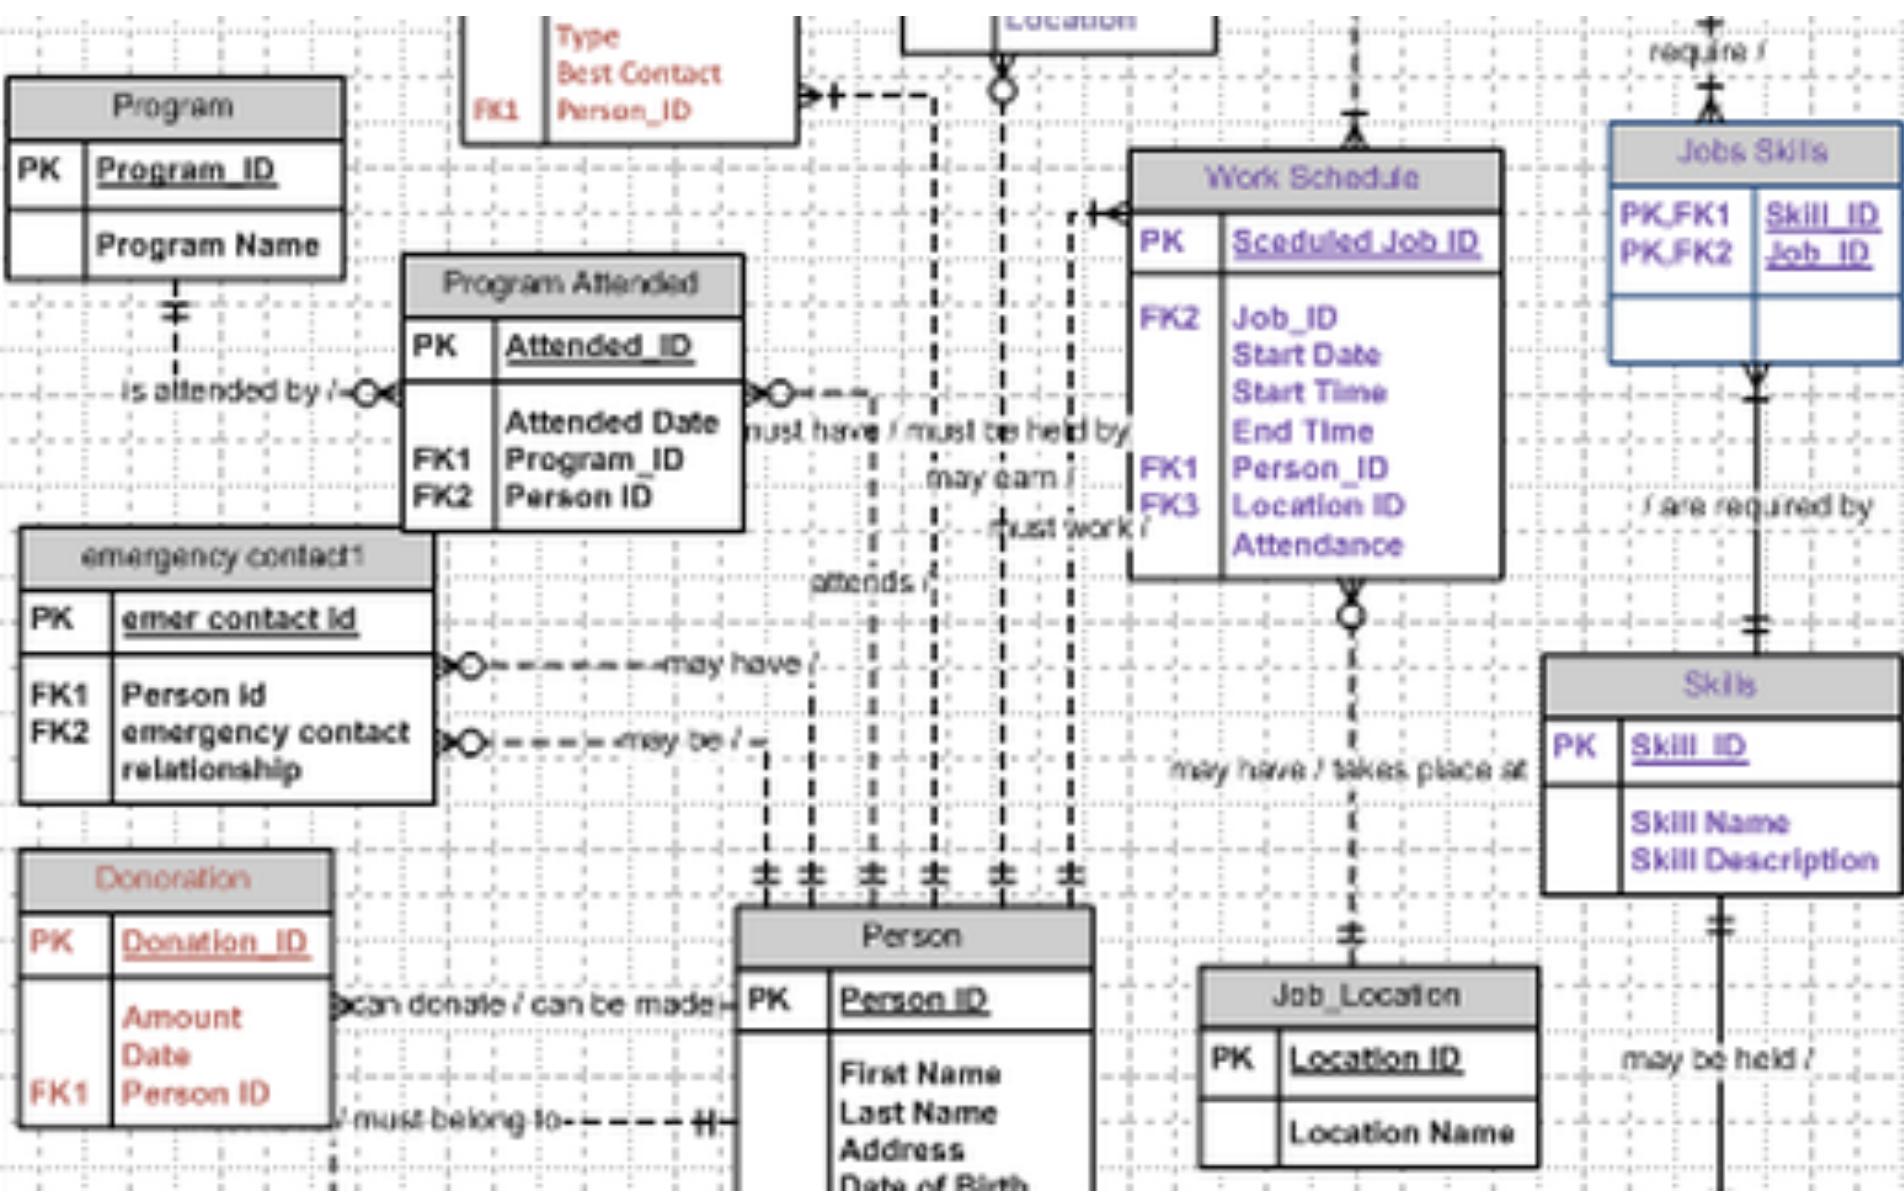
**MS Visio is a modeling tool for ER Diagrams**



Assumption: Class only offered once a day, 2 or more instructors per class possible.







Locations table that contains specific address of a specific...

### LOCATIONS

LOCATION_ID	NUMBER
STREET_ADDRESS	VARCHAR2
POSTAL_CODE	VARCHAR2
CITY	VARCHAR2
STATE_PROVINCE	VARCHAR2
COUNTRY_ID	CHAR

+ Constraints  
+ Indexes

W LOC\_C\_ID\_FK

countrytable. Contains 25 rows. References w...

### COUNTRIES

COUNTRY_ID	CHAR
COUNTRY_NAME	VARCHAR2
REGION_ID	NUMBER

+ Constraints  
+ Indexes

W COUNTR\_REG\_FK

Regions table that contains re...

### REGIONS

Departments table that shows details of departments where employees work. Contains 27 rows; references w...

### DEPARTMENTS

DEPARTMENT_ID	NUMBER
DEPARTMENT_NAME	VARCHAR2
MANAGER_ID	NUMBER
DEPT_LOC_FK	NUMBER

+ Constraints  
+ Indexes

EMP\_DEPT\_FK

employees table. Contains 107 rows. References with de...

### Procedures

ADD_JOB_HIST...
P_EMP_ID
P_START_DATE
P_END_DATE
P_JOB_ID
P_DEPARTMENT_ID

SECURE\_DML

Database Diagram in dbForge Studio for Oracle

<http://www.devart.com/dbforge/oracle/studio/oracle-databases>

Table that stores job history of the employees. If an employee...

### JOB\_HISTORY

EMPLOYEE_ID	NUMBER
START_DATE	DATE
END_DATE	DATE
JOB_ID	NUMBER

+ Constraints  
+ Indexes

jobs table with job titles and salary ranges. Contains 19 r...

### JOB

JOB_ID	NUMBER
JOB_TITLE	VARCHAR2
MIN_SALARY	NUMBER

+ Constraints  
+ Indexes

### EMP\_DETAILS\_VIEW

EMPLOYEE_ID	NUMBER
JOB_ID	VARCHAR2
MANAGER_ID	NUMBER
DEPARTMENT_ID	NUMBER
LOCATION_ID	NUMBER
COUNTRY_ID	CHAR
FIRST_NAME	VARCHAR2
LAST_NAME	VARCHAR2
SALARY	NUMBER

### Devert

#### HUMAN\_RESOURCES

Version:	3.50
Author:	Devert Company
Date:	6/19/2013
Copyrights:	© Devert. All rights reserved.

# E-R Diagrams

---

- ◆ There are variations to ER Diagrams
- ◆ For this course we are focused on the ones we discussed. This is based on the diagrams that are usually auto-produced for databases.
  - Designate keys
  - Attributes inside box (*not ovals outside of it*)
- ◆ You need to be able to read an ER Diagram as you may have one to work with for coding

---

# Structured Query Language

# **SQL**

# Select Basics

- ◆ **SQL** is a language to write queries over the data
- ◆ Most basic of all SQL statements

**SELECT *columnlist* FROM *table***

- *Columnlist* named by comma separated list.

**SELECT firstName, lastName FROM t\_contacts**

- Use \* as a shortcut for all columns.

**SELECT \* FROM t\_contacts**

All caps for  
SQL keywords  
is a common  
convention

- **Aliasing** can be used to change the display

- This is useful to improve readability of reports

**SELECT data3 AS start\_date FROM t\_contacts**

# Filtering and Ordering

---

```
SELECT columnlist  
FROM table  
WHERE columnCondition  
ORDER BY columnOrder [ASC/DESC]
```

- ◆ **WHERE** clause filters based on columns using Boolean and logical operators

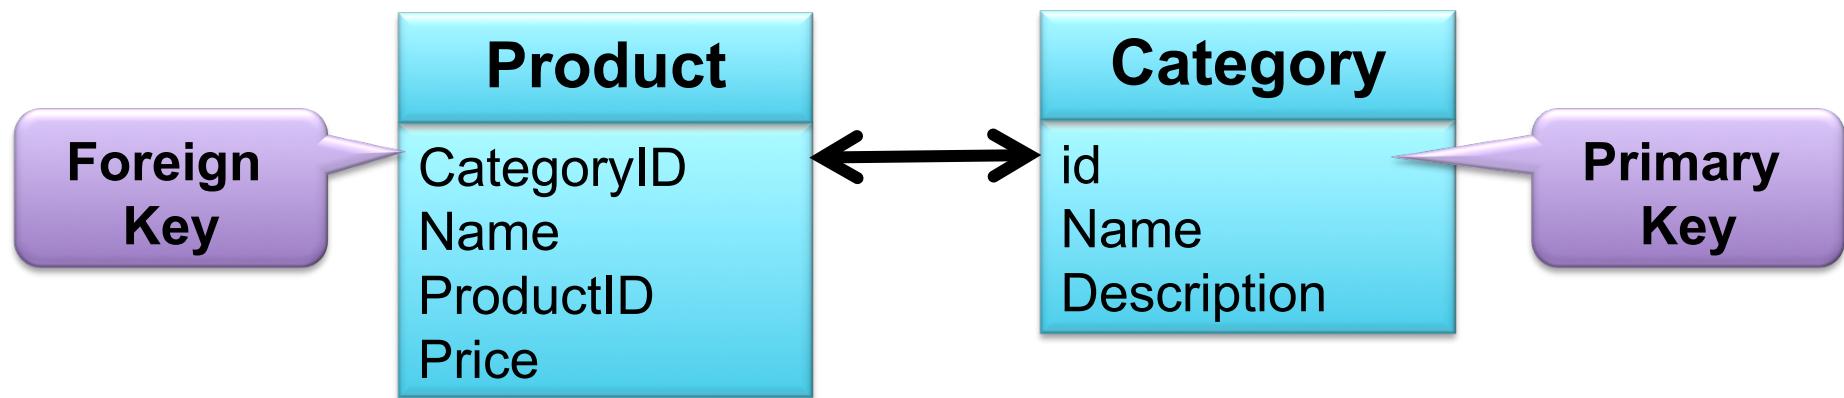
```
SELECT * FROM t_contacts  
WHERE lastName = "Boese"
```

- ◆ **ORDER BY** defaults to **ASC**

```
SELECT * FROM t_contacts  
ORDER BY lastName DESC
```

# Joining tables

- ◆ Relationships are generally modeled through “key” columns.
- ◆ Primary and Foreign keys
- ◆ Example – “CategoryID” column in the Product table references the “id” column in the Category table.



# Join Syntax

---

- ◆ Get results from multiple tables using where

```
SELECT columnlist  
FROM table1, table2  
WHERE table1.col_1 = table2.col_2
```

- ◆ Example

```
SELECT name, price, description  
FROM Product, Category  
WHERE Category.id = Product.categoryID
```

# Join Syntax

---

- ◆ Get results from multiple tables

```
SELECT columnlist  
FROM table1  
JOIN table2 ON table1.col_1 = table2.col_2
```

- ◆ Example

```
SELECT name, price, description  
FROM Product  
JOIN Category ON categoryID = id
```

# Join Syntax

---

- ◆ WHERE vs JOIN?
- ◆ SIDEBAR... *advanced concepts [optional]*

*INNER JOIN is ANSI syntax which you should use.*

*It is generally considered more readable, especially when you join lots of tables.  
It can also be easily replaced with an OUTER JOIN whenever a need arises.*

*The WHERE syntax is more relational model oriented.*

*A result of two tables JOIN'ed is a cartesian product of the tables to which a filter is applied which selects only those rows with joining columns matching.*

*It's easier to see this with the WHERE syntax.*

*Also note that MySQL also has a STRAIGHT\_JOIN clause.*

*Using this clause, you can control the JOIN order: which table is scanned in the outer loop and which one is in the inner loop.*

*You cannot control this in MySQL using WHERE syntax.*

# Name Resolution

---

```
SELECT name, price, description  
FROM Product  
JOIN Category ON categoryId = id
```

- ◆ Query engine must be able to resolve columns;  
in the example: “name” is ambiguous if there are columns  
called “name” in both tables.
- ◆ Fully qualified syntax can be used

**SELECT** tableName.columnName, ...

Ex:

**SELECT** Product.Name, Category.Name, ...

# Aliases

---

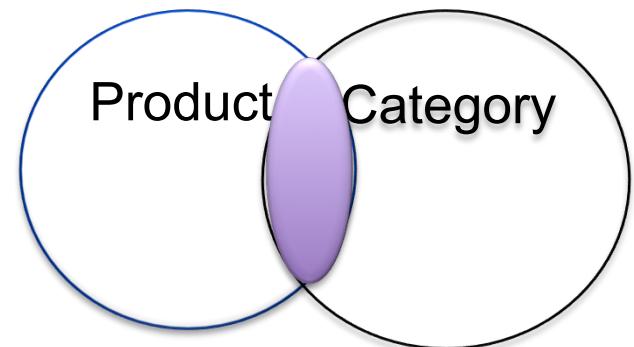
## Aliases

- Shorten typing
  - Rename columns.
- ◆ On tables and/or column names
    - FROM Product p
  - ◆ “AS” keyword optional
    - (*usually used on column names but not on table names*)
    - *p.name as “ProductName”*

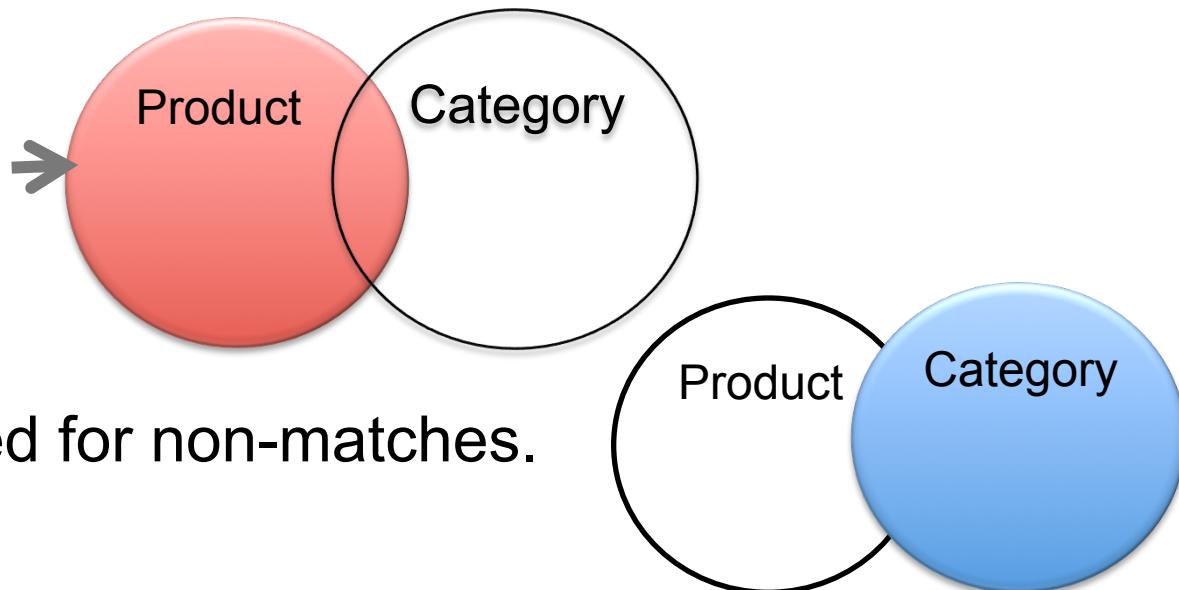
```
SELECT p.name as “ProductName”, c.name  
FROM Product p  
JOIN category c ON p.categoryID = c.id
```

# JOIN Types

- ◆ **JOIN** is shortcut for **INNER JOIN**
  - Only records that match are returned.



- ◆ **OUTER JOIN** may return records that do not match the filter condition.
  - **LEFT JOIN, RIGHT JOIN** syntax signifies which table has “optional” values.



- ◆ Null values returned for non-matches.

# Inner Join

```
SELECT *
FROM Customers C
INNER JOIN Orders O
ON O.CustomerId = C.CustomerId
```

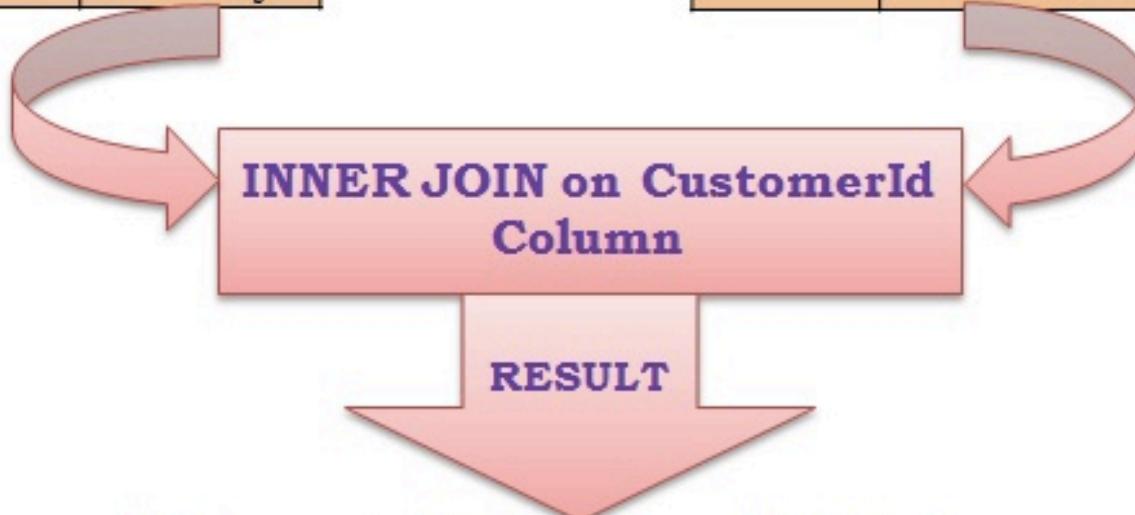
## INNER JOIN

Customers

CustomerId	Name
1	Shree
2	Kalpana
3	Basavaraj

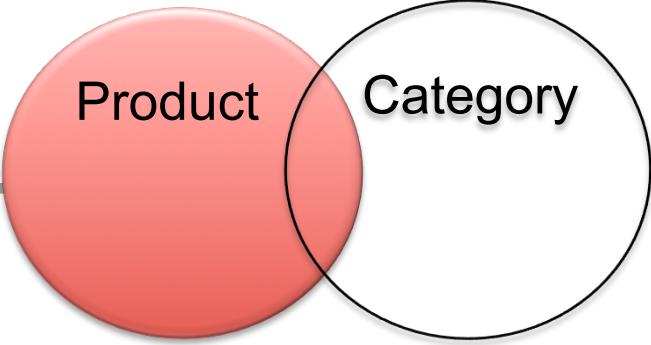
Orders

OrderId	CustomerId	OrderDate
100	1	2014-01-29 23:56:57.700
200	4	2014-01-30 23:56:57.700
300	3	2014-01-31 23:56:57.700



CustomerId	Name	OrderId	CustomerId	OrderDate
1	Shree	100	1	2014-01-30 23:48:32.850
3	Basavaraj	300	3	2014-02-01 23:48:32.853

# Left Join



- ◆ Left join, sometimes called left outer join
- ◆ The LEFT JOIN keyword returns all the rows from the left table (Product), even if there are no matches in the right table (Category).
- ◆ A *list of all products and their categories even if a product not associated to a category.*  
*If no products use a particular category, that category does not show up*

```
SELECT p.name, c.name  
FROM Product p  
LEFT JOIN category c ON p.categoryID = c.id
```

# Left Join

```
SELECT *
FROM Customers C
LEFT OUTER JOIN Orders O
ON O.CustomerId = C.CustomerId
```

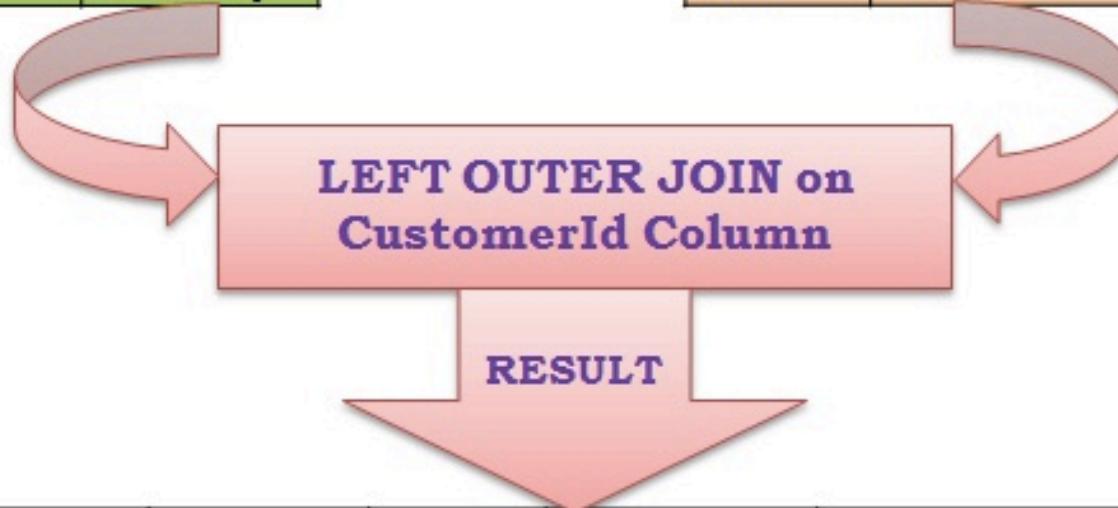
## LEFT OUTER JOIN

Customers

CustomerId	Name
1	Shree
2	Kalpana
3	Basavaraj

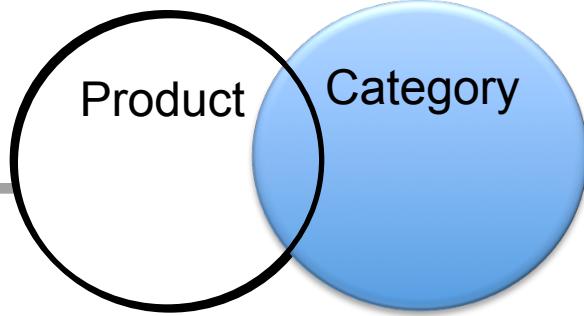
Orders

OrderId	CustomerId	OrderDate
100	1	2014-01-29 23:56:57.700
200	4	2014-01-30 23:56:57.700
300	3	2014-01-31 23:56:57.700



CustomerId	Name	OrderId	CustomerId	OrderDate
1	Shree	100	1	2014-01-30 23:48:32.850
2	Kalpana	NULL	NULL	NULL
3	Basavaraj	300	3	2014-02-01 23:48:32.853

# Right Join



- ◆ Right join, sometimes called right outer join
- ◆ The **RIGHT JOIN** keyword returns all the rows from the right table (**Category**), even if there are no matches in the left table (**Product**).
- ◆ *A list of all categories and their matching products even if a category is not associated to a product.*  
*If no categories reference a particular product, that product does not show up*

```
SELECT p.name, c.name  
FROM Product p  
RIGHT JOIN category c ON p.categoryID = c.id
```

# Right Outer Join

```
SELECT *  
FROM Customers C  
      RIGHT OUTER JOIN Orders O  
    ON O.CustomerId = C.CustomerId
```

## RIGHT OUTER JOIN

Customers

CustomerId	Name
1	Shree
2	Kalpana
3	Basavaraj

Orders

OrderId	CustomerId	OrderDate
100	1	2014-01-29 23:56:57.700
200	4	2014-01-30 23:56:57.700
300	3	2014-01-31 23:56:57.700

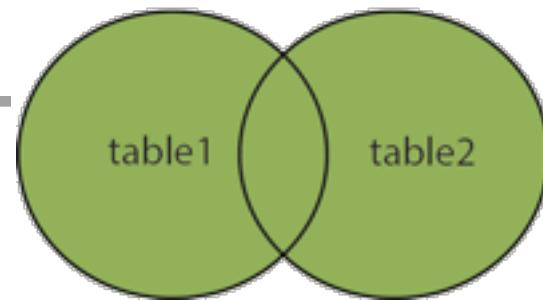
RIGHT OUTER JOIN on  
CustomerId Column

RESULT

CustomerId	Name	OrderId	CustomerId	OrderDate
1	Shree	100	1	2014-01-30 23:48:32.850
NULL	NULL	200	4	2014-01-31 23:48:32.853
3	Basavaraj	300	3	2014-02-01 00:48:32.853

FULL OUTER JOIN

# Full Outer Join



- ◆ Full outer join
- ◆ The FULL OUTER JOIN keyword returns all the rows from both tables.

```
SELECT p.name, c.name  
FROM Product p  
FULL OUTER JOIN category c ON p.categoryID = c.id
```

```

SELECT *
FROM Customers C
FULL OUTER JOIN Orders O
ON O.CustomerId = C.CustomerId

```

# Full Outer Join

## FULL OUTER JOIN

**Customers**

CustomerId	Name
1	Shree
2	Kalpana
3	Basavaraj

**Orders**

OrderId	CustomerId	OrderDate
100	1	2014-01-29 23:56:57.700
200	4	2014-01-30 23:56:57.700
300	3	2014-01-31 23:56:57.700

FULL OUTER JOIN on  
CustomerId Column

RESULT

CustomerId	Name	OrderId	CustomerId	OrderDate
1	Shree	100	1	2014-01-30 23:48:32.850
2	Kalpana	NULL	NULL	NULL
3	Basavaraj	300	3	2014-02-01 23:48:32.853
NULL	NULL	200	4	2014-01-31 23:48:32.853

# Aggregate Functions

---

- ◆ Return a single value calculated from values in the column.
- ◆ Examples: **AVG()**, **COUNT()**, **MAX()**, **MIN()**, **SUM()**
- ◆ Can use a “**GROUP BY**” clause that includes all columns not aggregated to achieve results by groups

```
SELECT AVG(p.Price) 'Average Price'  
, c.Name 'Category Name'  
, c.Description 'Category Description'  
FROM Product p JOIN Category c  
ON p.categoryID = c.idCategory  
GROUP BY c.Name, c.Description
```

# Unions

---

- ◆ Sometimes easier to break a query into pieces.
- ◆ Combine the separate result sets into one final set.
  - Default is to remove duplicates
  - Include duplicates with: UNION ALL
- ◆ Column number and type must match

```
SELECT column_name(s) FROM table1  
UNION  
SELECT column_name(s) FROM table2;
```

# Modifying Data

---

- ◆ Update statement used to change existing data.

```
UPDATE table_name  
SET column1 = value1, column2 = value2,...  
WHERE some_column = some_value;
```

- ◆ Delete used to remove records

```
DELETE FROM table_name  
WHERE some_column = some_value;
```

- ◆ Insert used to create new records

```
INSERT INTO table_name (column1, column2, ...)  
VALUES (value1, value2, ...);
```

# Transactions

---

## ◆ ACID Transactions

- **Atomic** – transactions are all or nothing
  - **Consistent** – all transactions must leave database in a valid state as per constraints
  - **Isolated** – ordering is preserved in face of concurrency
  - **Durable** – committed transaction persists even in face of system failure
- ◆ Multiple statements are either completely executed or fail as a whole.
- Ensures integrity of the database.

---

# **STORED PROCEDURES**

# Stored Procedures

---

- ◆ Pre-compiled SQL statement that reside in the database.

```
CREATE PROCEDURE uspGetAddress  
AS  
SELECT * FROM AdventureWorks.Person.Address  
GO
```

- ◆ Usually consist of input parameters that may perform several related actions.
- ◆ Called from other statements using EXEC.

```
EXEC uspGetAddress
```

# Programming Frameworks

---

- ◆ Allow for interacting with databases via language constructs.
- ◆ Connections, statements, and result sets are represented by classes.
- ◆ Can send dynamically generated SQL to the database and process the results.
- ◆ Example: Yii framework (*see slides from previous week*)

# Object Relational Mapping

---

- ◆ “**Domain object modeling**” implies a class represents a single table.
  - Store table represented by a Store class.
- ◆ Manually mapping columns to object properties.
- ◆ Modeling objects via annotations.
  - Queries done in programming language.

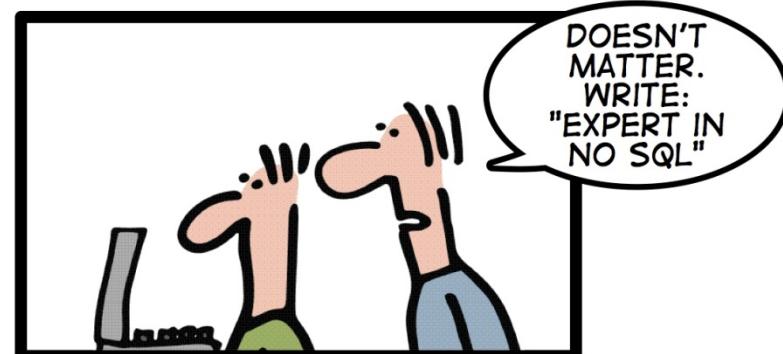
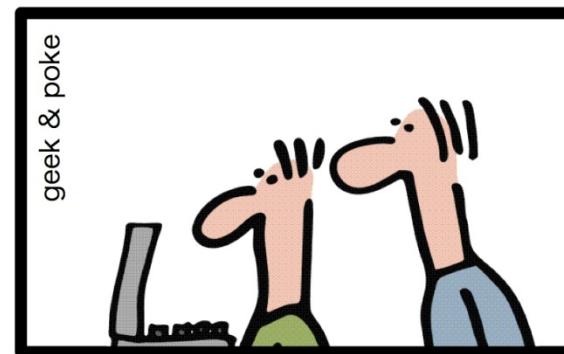
Example: **ActiveRecord** in Ruby on Rails & many Frameworks

# HOW TO WRITE A CV

---

**NOSQL: NO SQL**

**NOSQL:**  
**NOT ONLY SQL**



Leverage the NoSQL boom

---

## NoSQL:

A database that provides a mechanism for storage and retrieval of data that is modeled in means **other than the tabular relations** used in relational databases. Motivations for this approach include **simplicity of design, horizontal scaling** and finer **control over availability**.

-[Wikipedia](#)

# Use Cases

---

## Motivations

- ◆ Big data
- ◆ Scalability
- ◆ Data format
- ◆ Manageability

## Use Cases

- Massive write performance.
- Fast key value look ups.
- Flexible schema and data types.
- No single point of failure.
- Fast prototyping and development.
- Out of the box scalability.
- Easy maintenance.

# Big Data

---

- ◆ Collect
- ◆ Store
- ◆ Organize
- ◆ Analyze
- ◆ Share

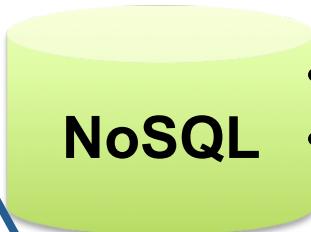
Data growth outruns the ability to manage it so we need **scalable** solutions.

# RDBMS vs NoSQL

- Oracle
- DB2
- SQL Server
- MySQL



**Availability**



- Cassandra
- CouchDB

**Consistency**

NoSQL

- Big Table
- MongoDB
- Redis

**Partition  
Tolerance**

---

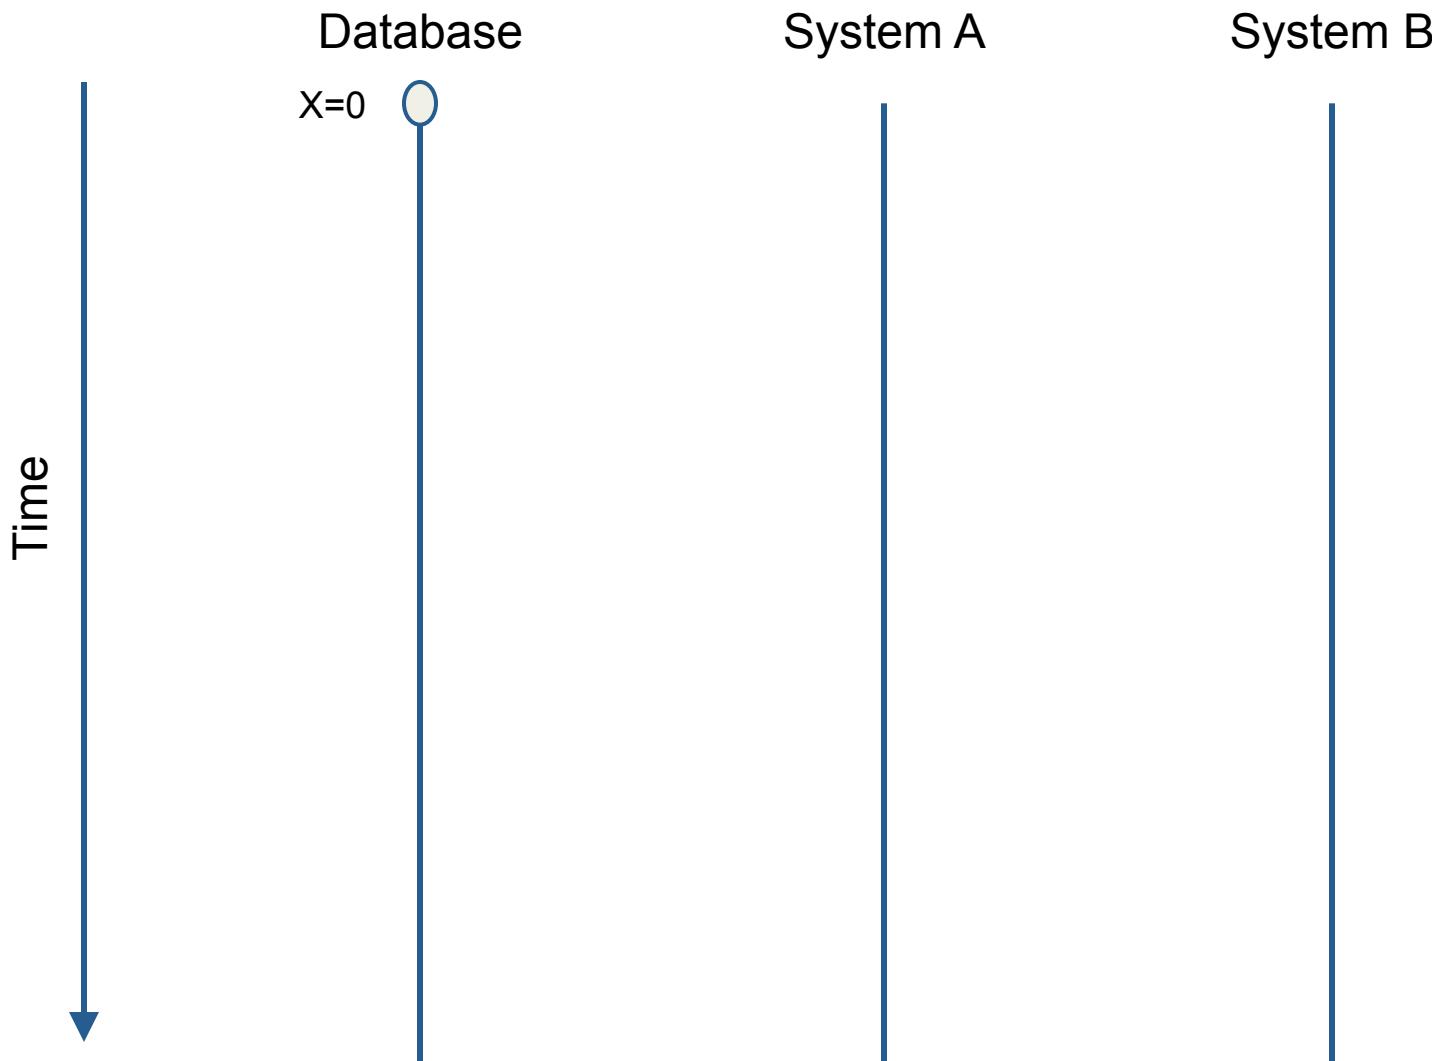
- ◆ NoSQL DBs (often) Lack:

- SQL-based Interfaces
- ACID Transactions
- Strong Consistency

*By giving up ACID properties, one can achieve higher performance and scalability.*

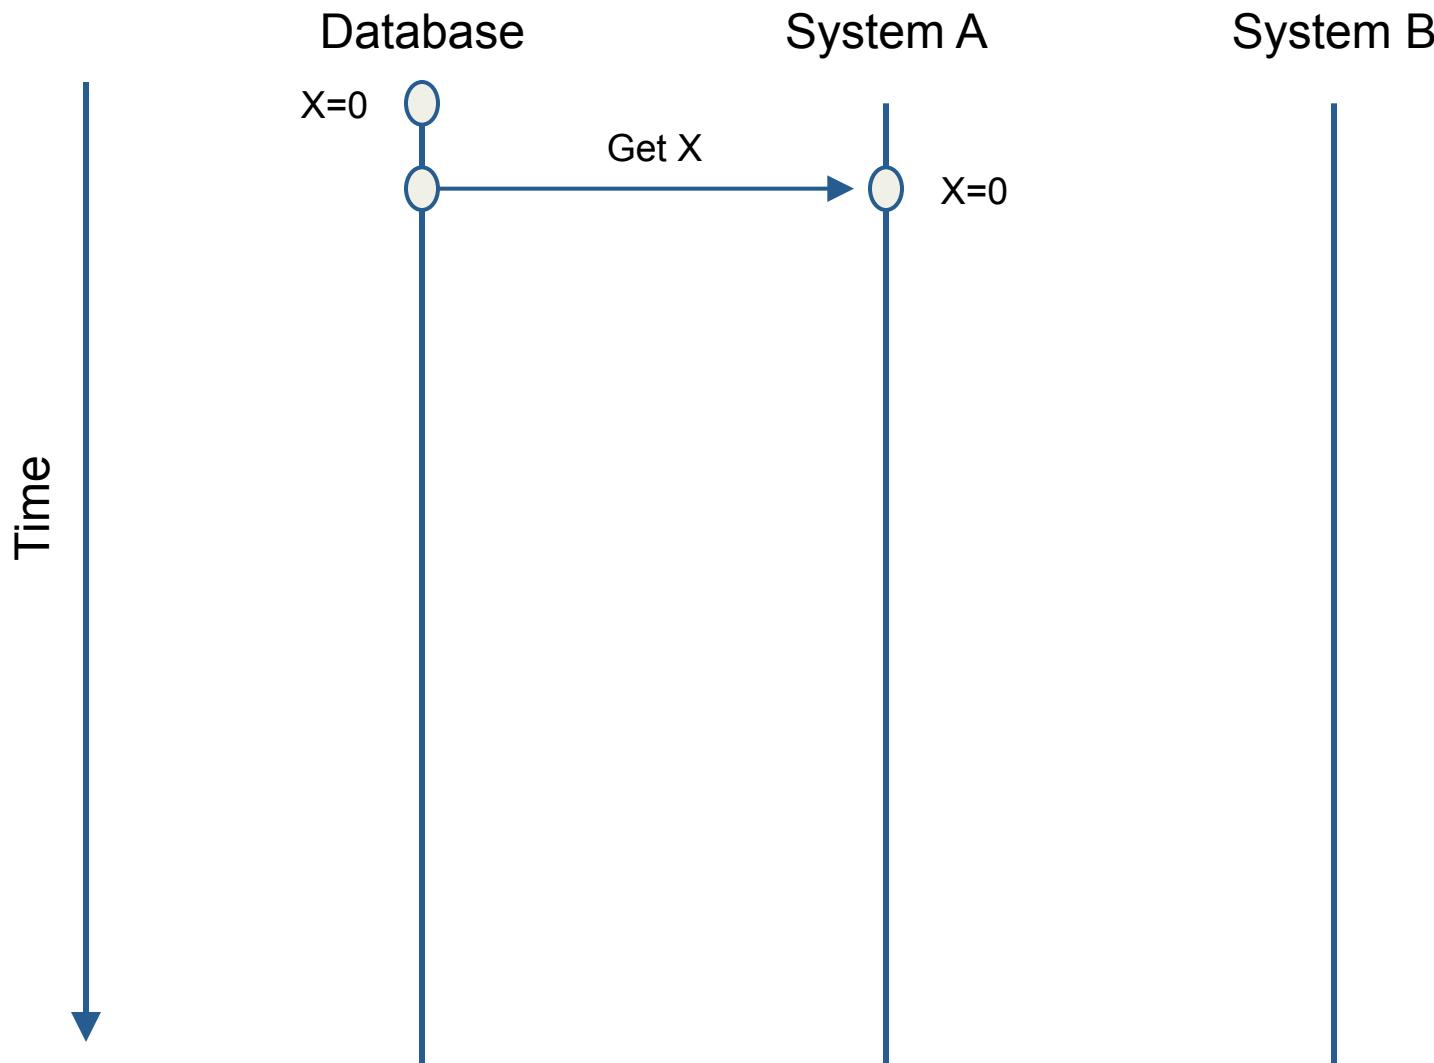
# Strong Consistency

---

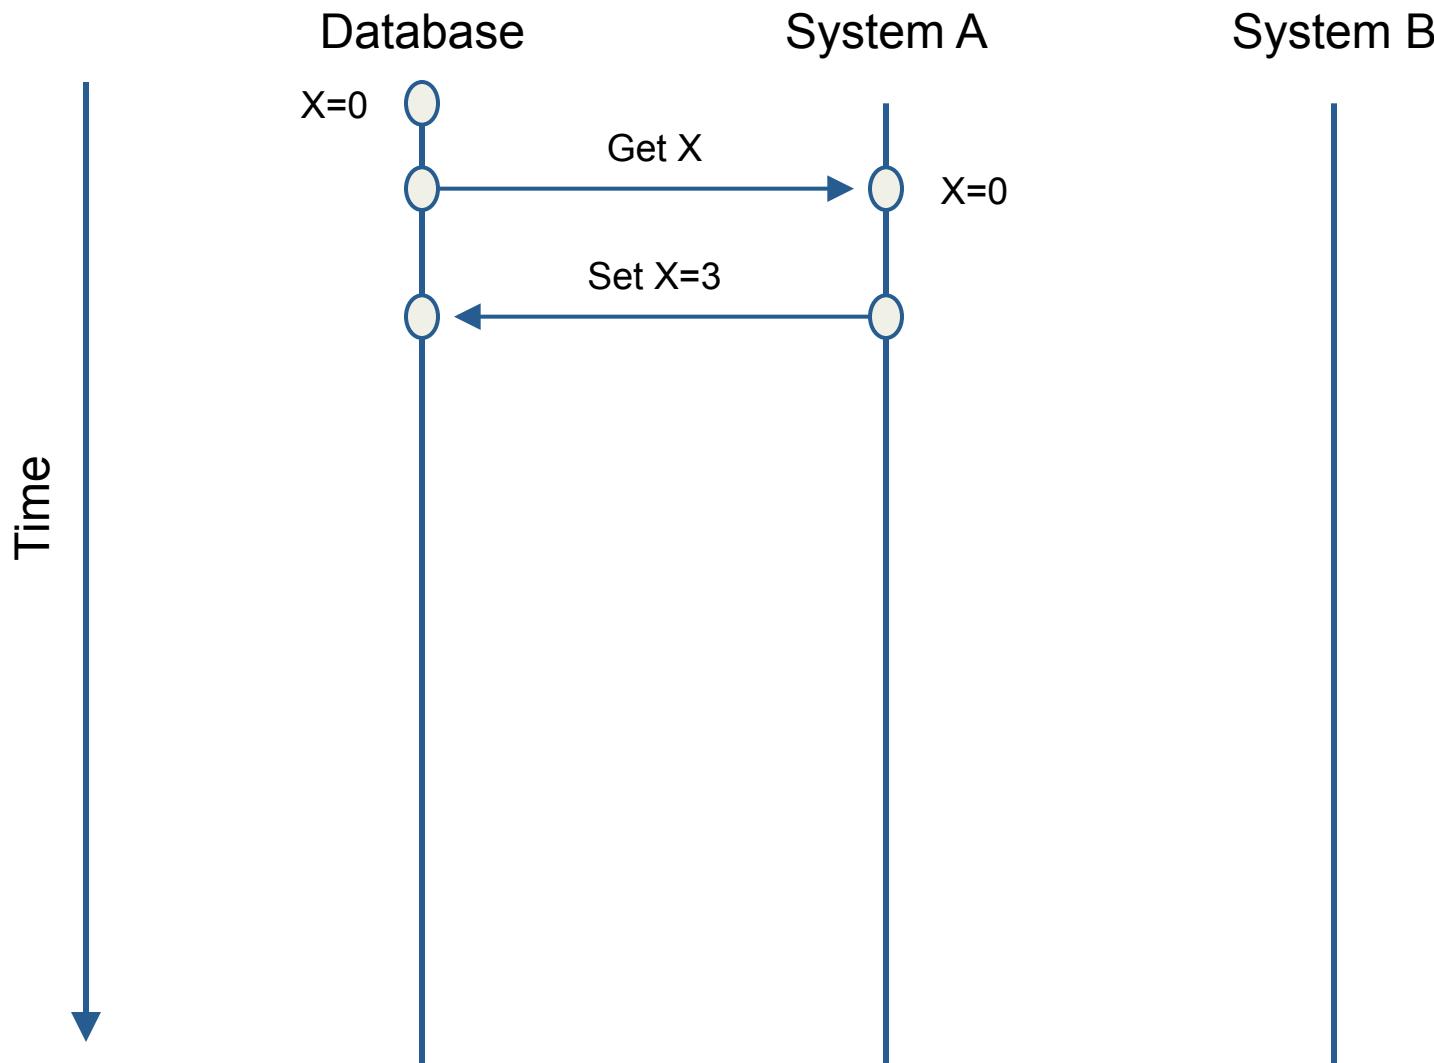


# Strong Consistency

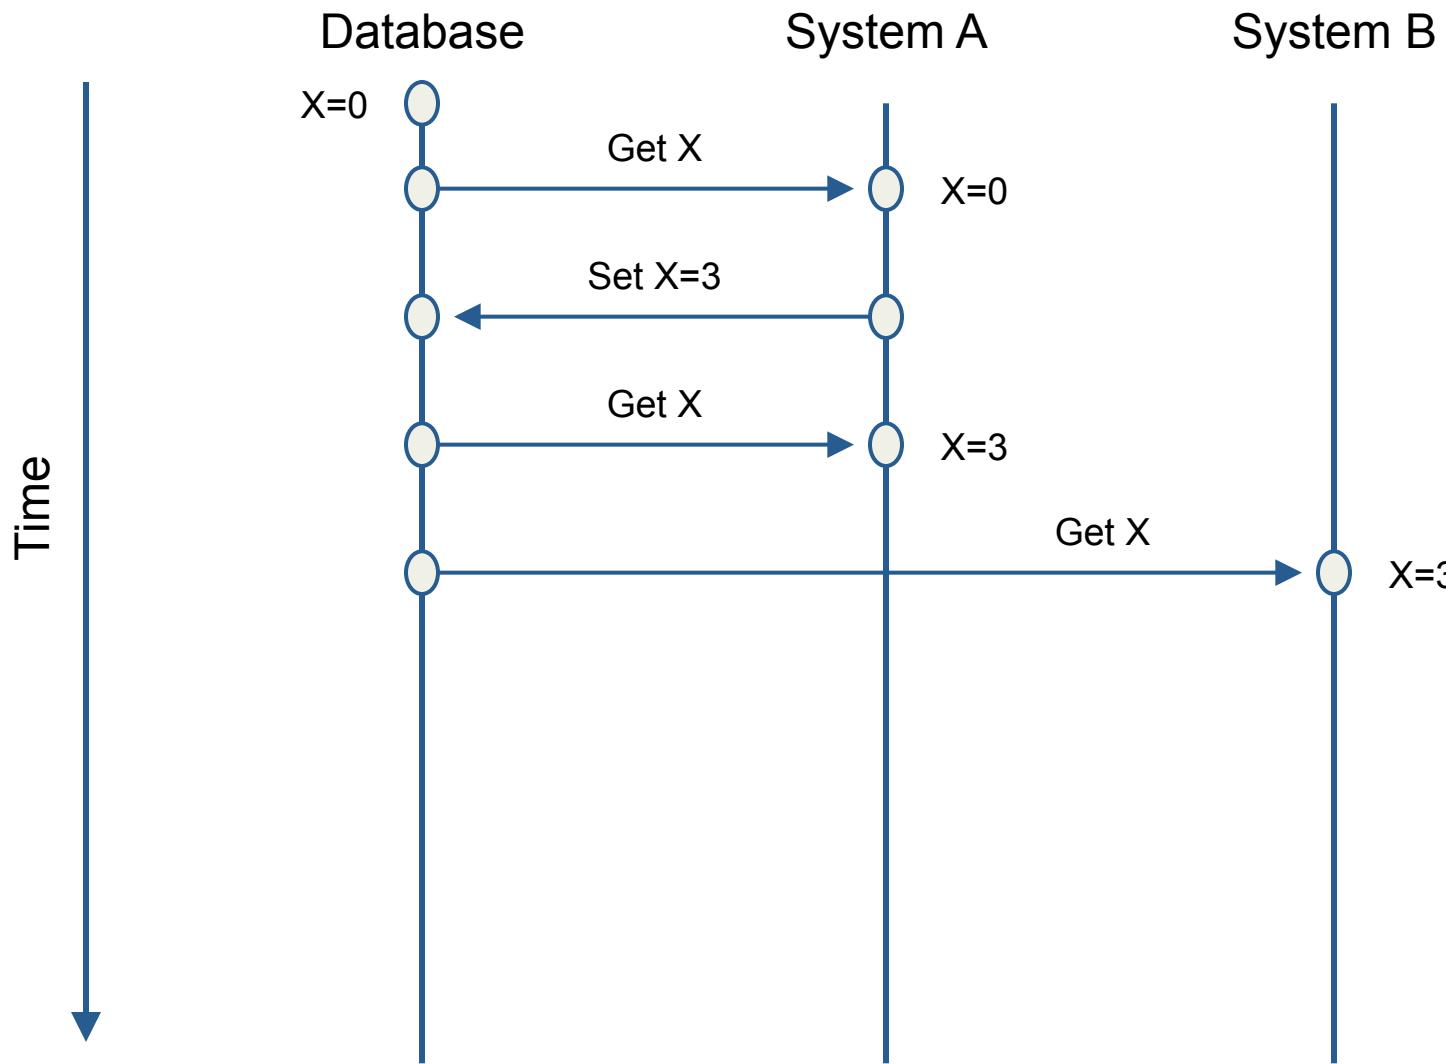
---



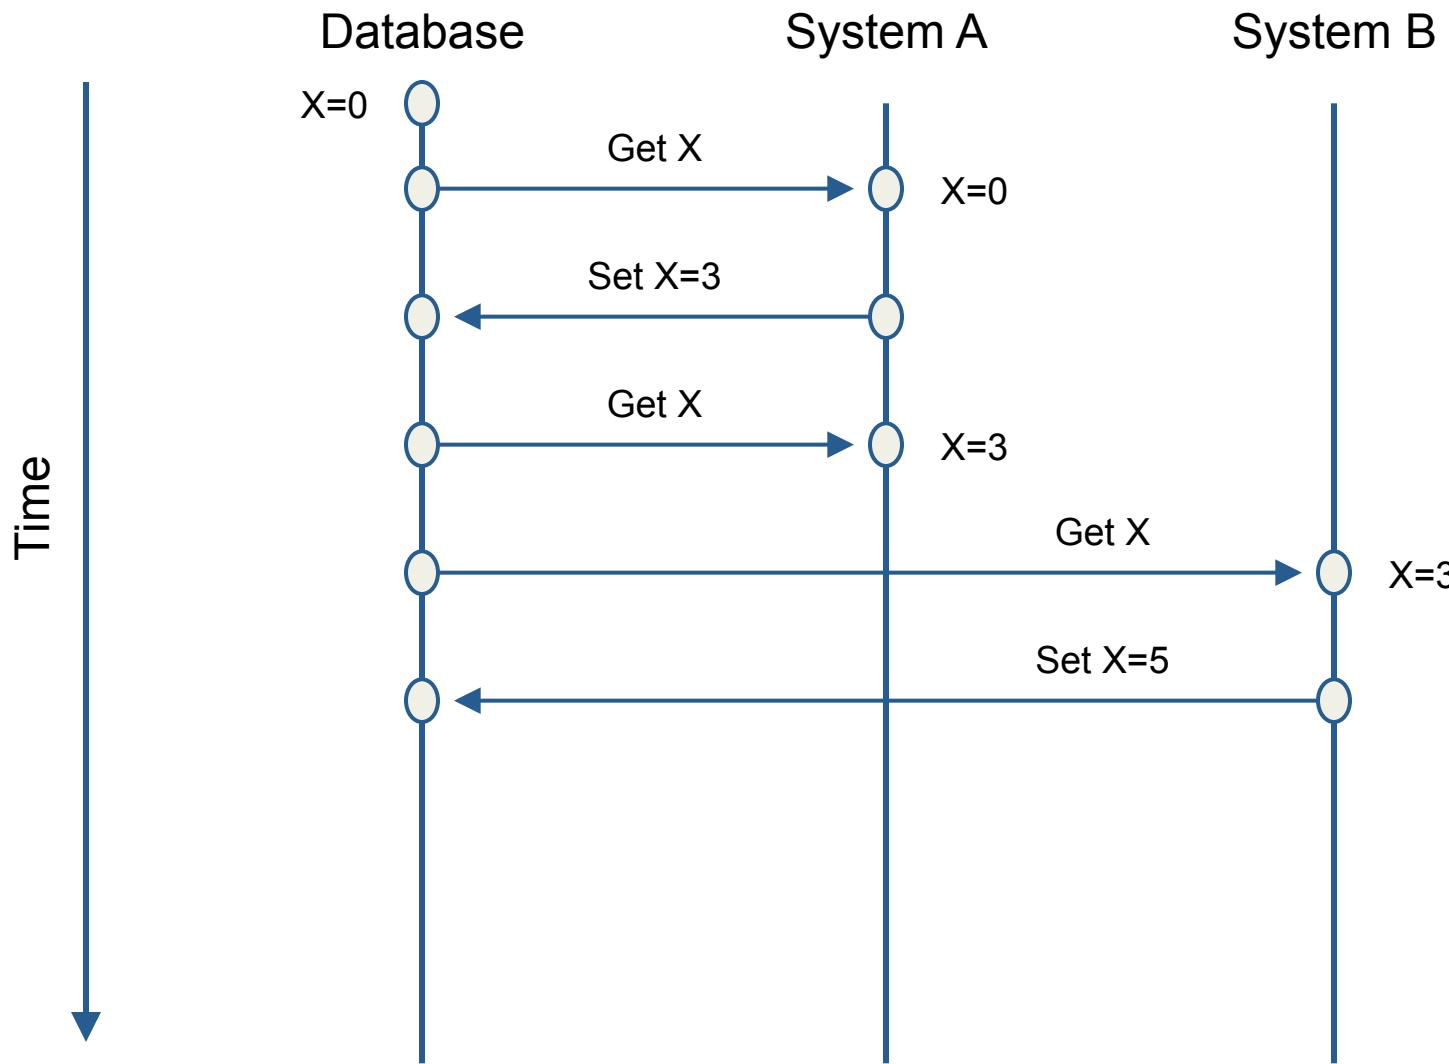
# Strong Consistency



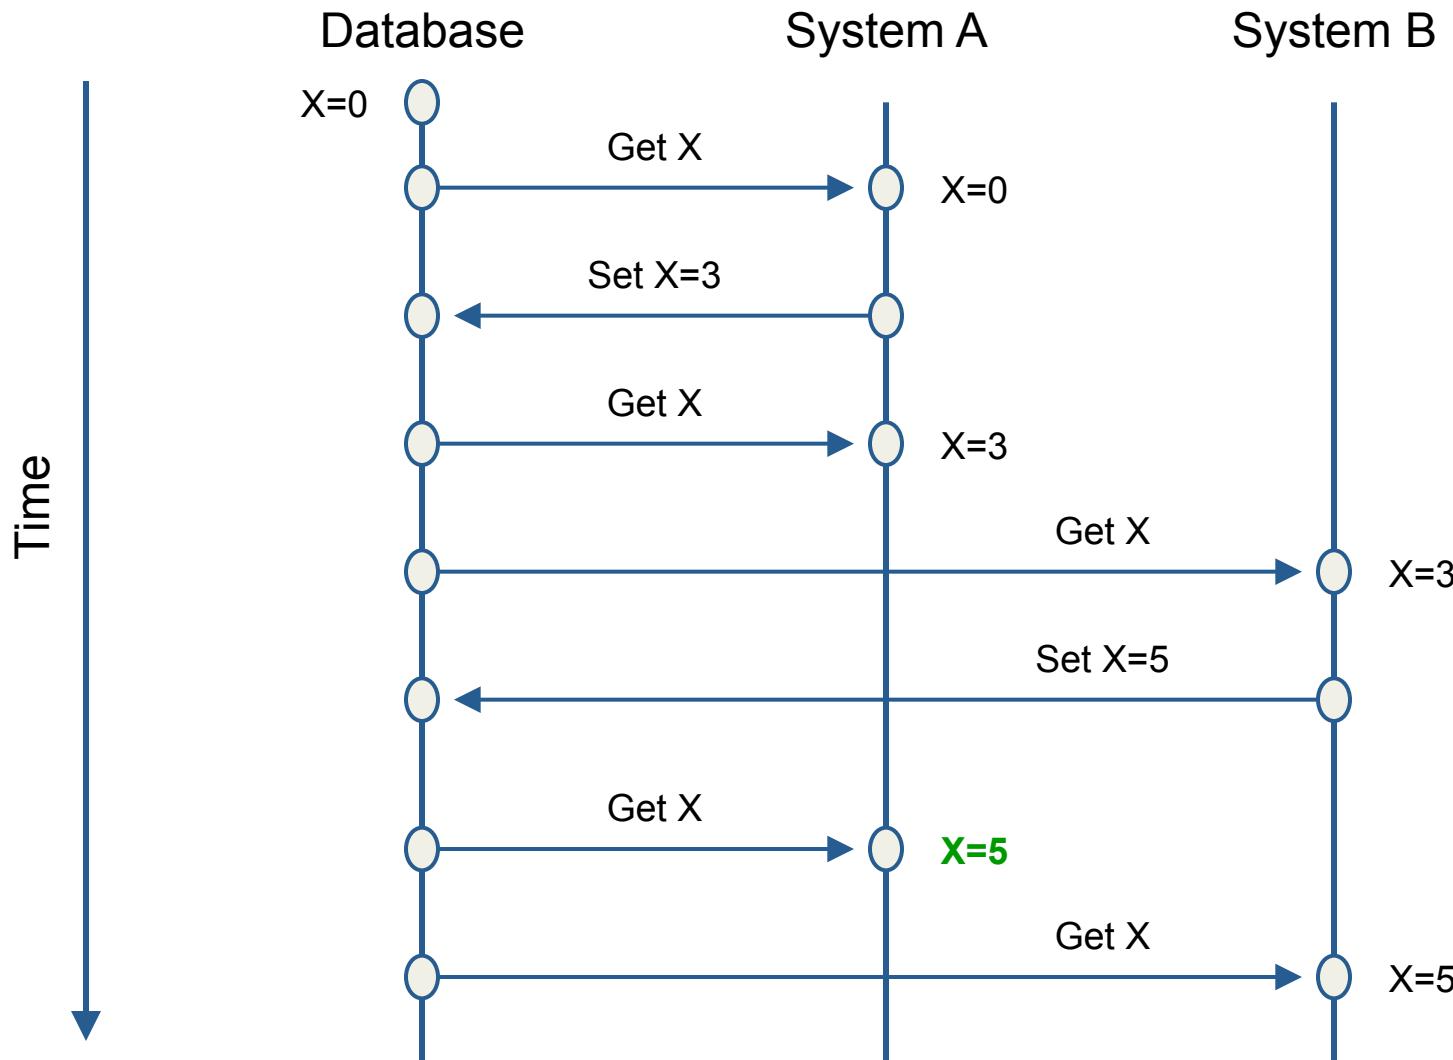
# Strong Consistency



# Strong Consistency

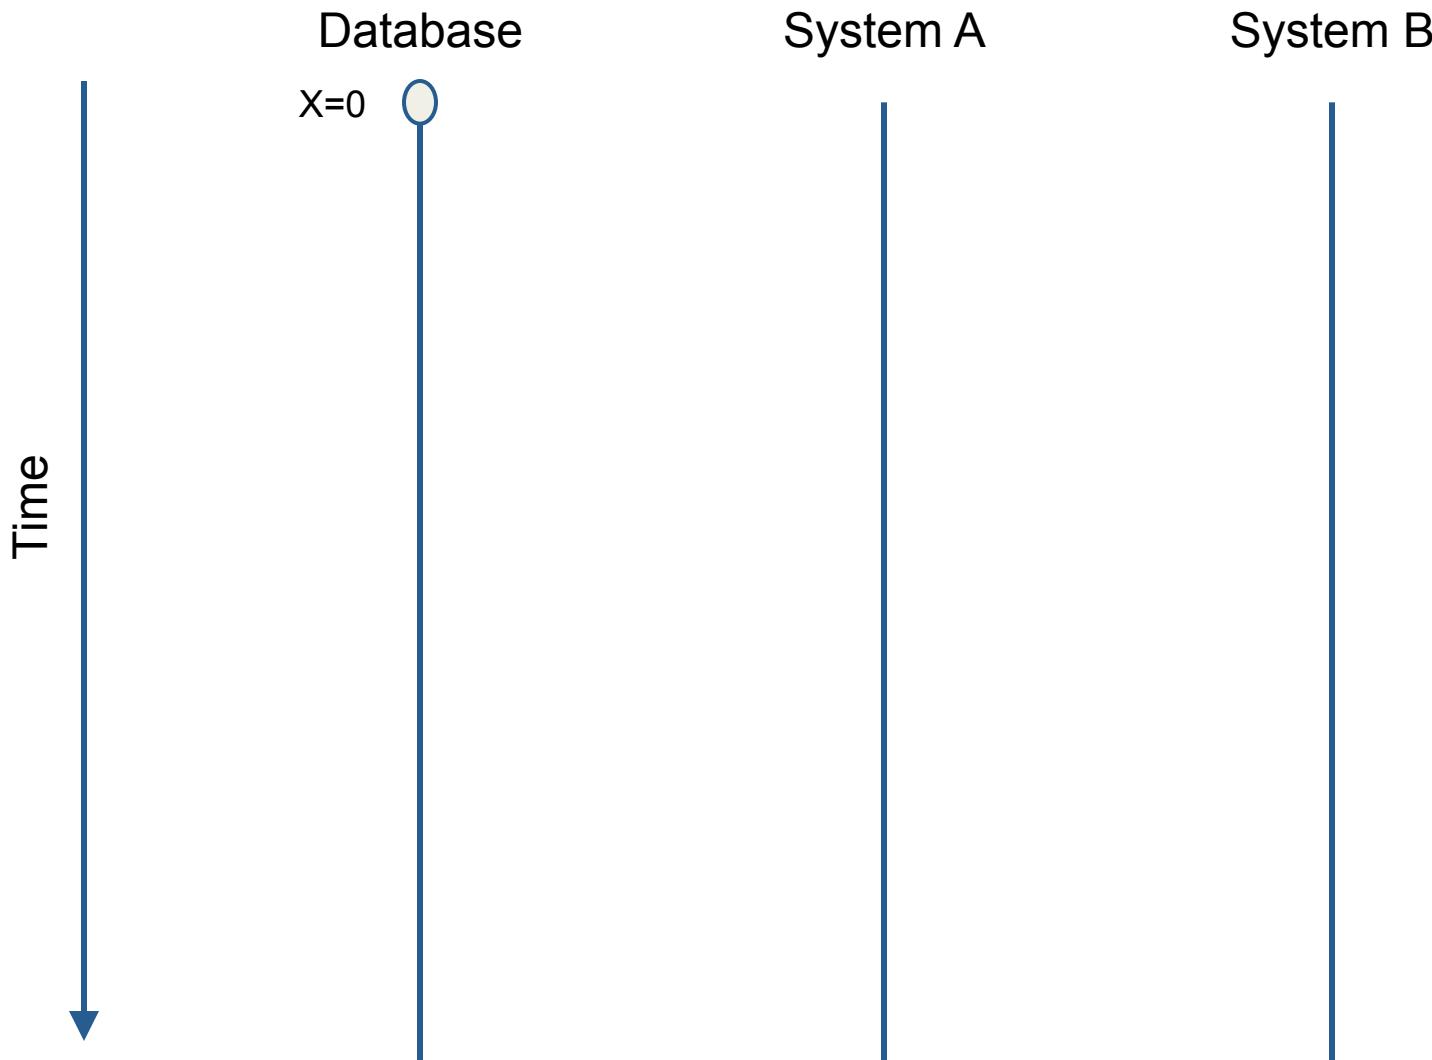


# Strong Consistency



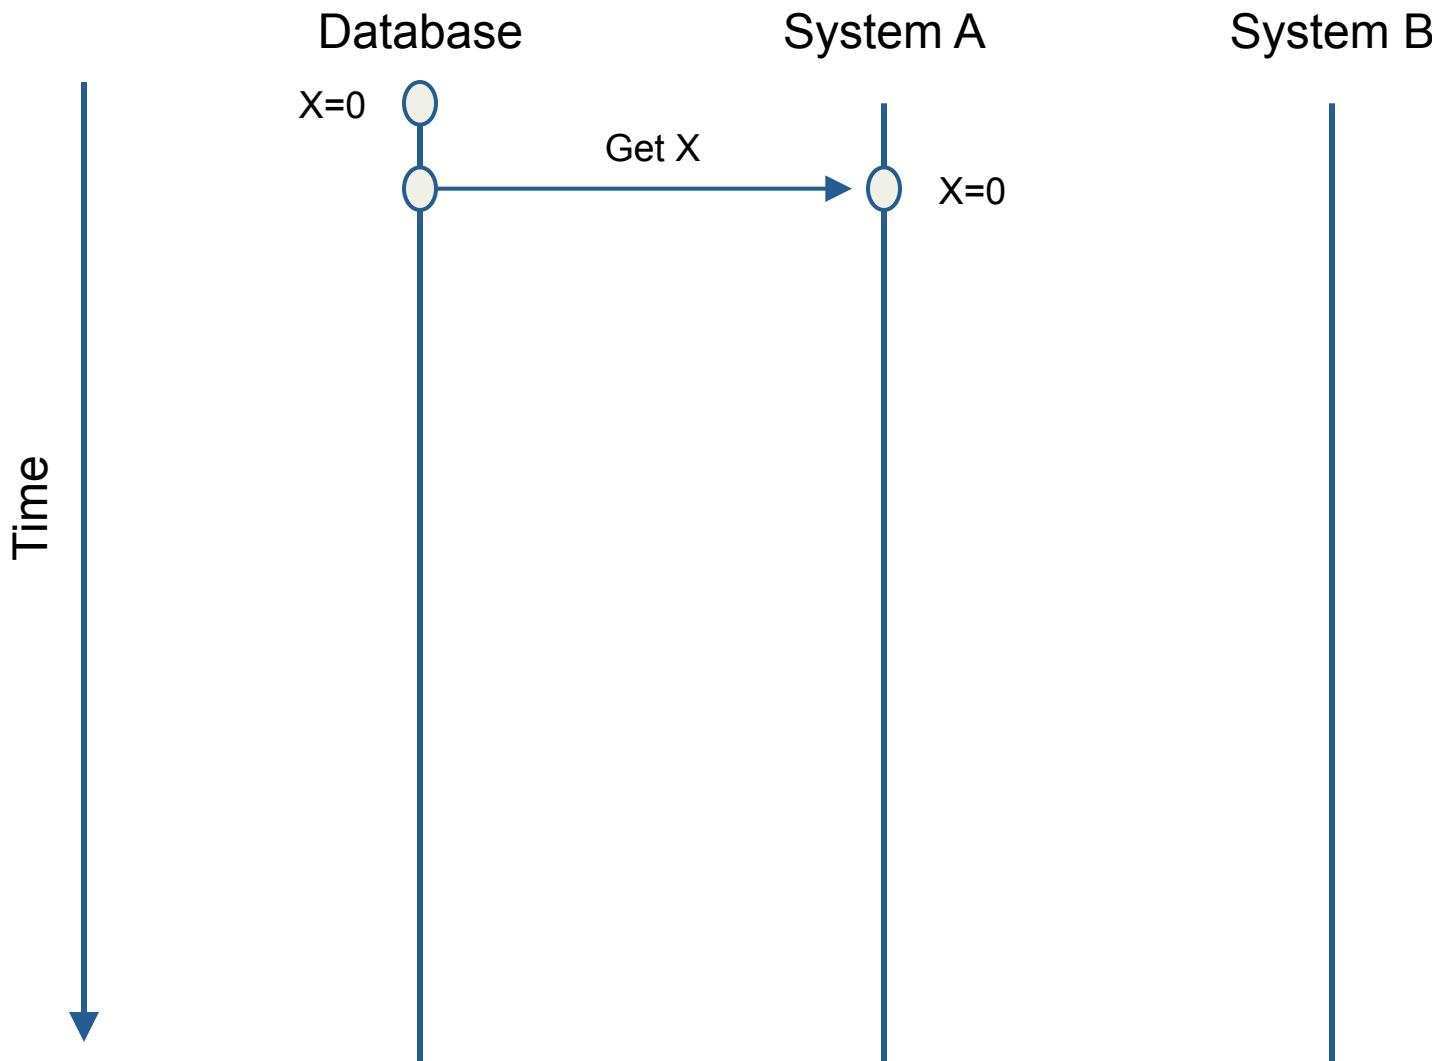
# Weak Consistency

---

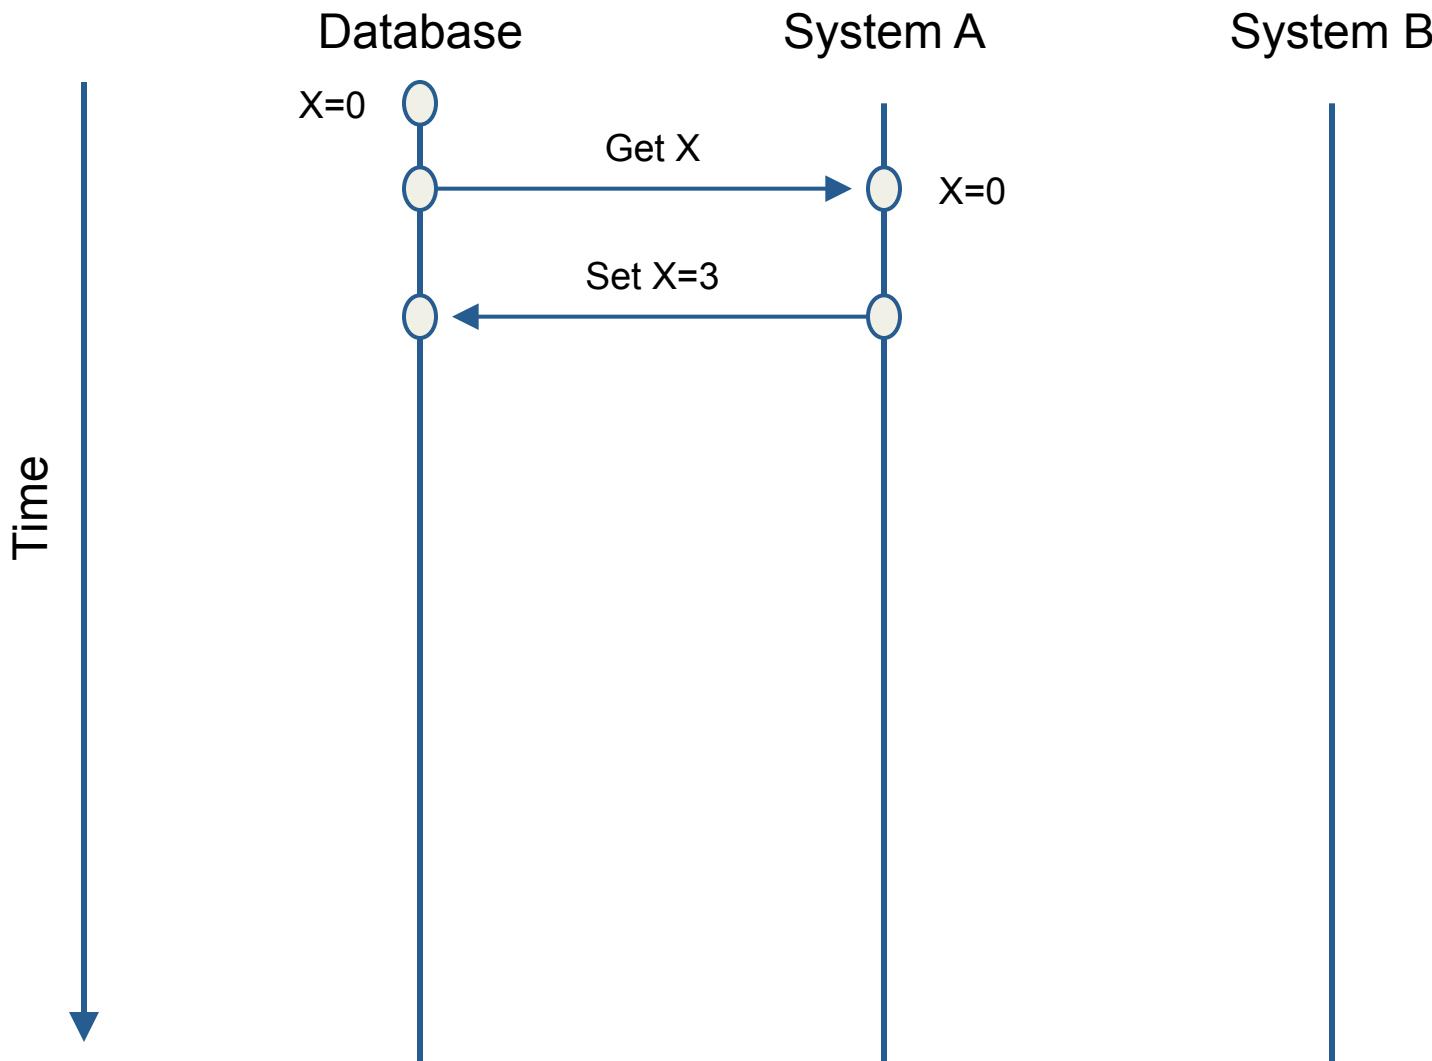


# Weak Consistency

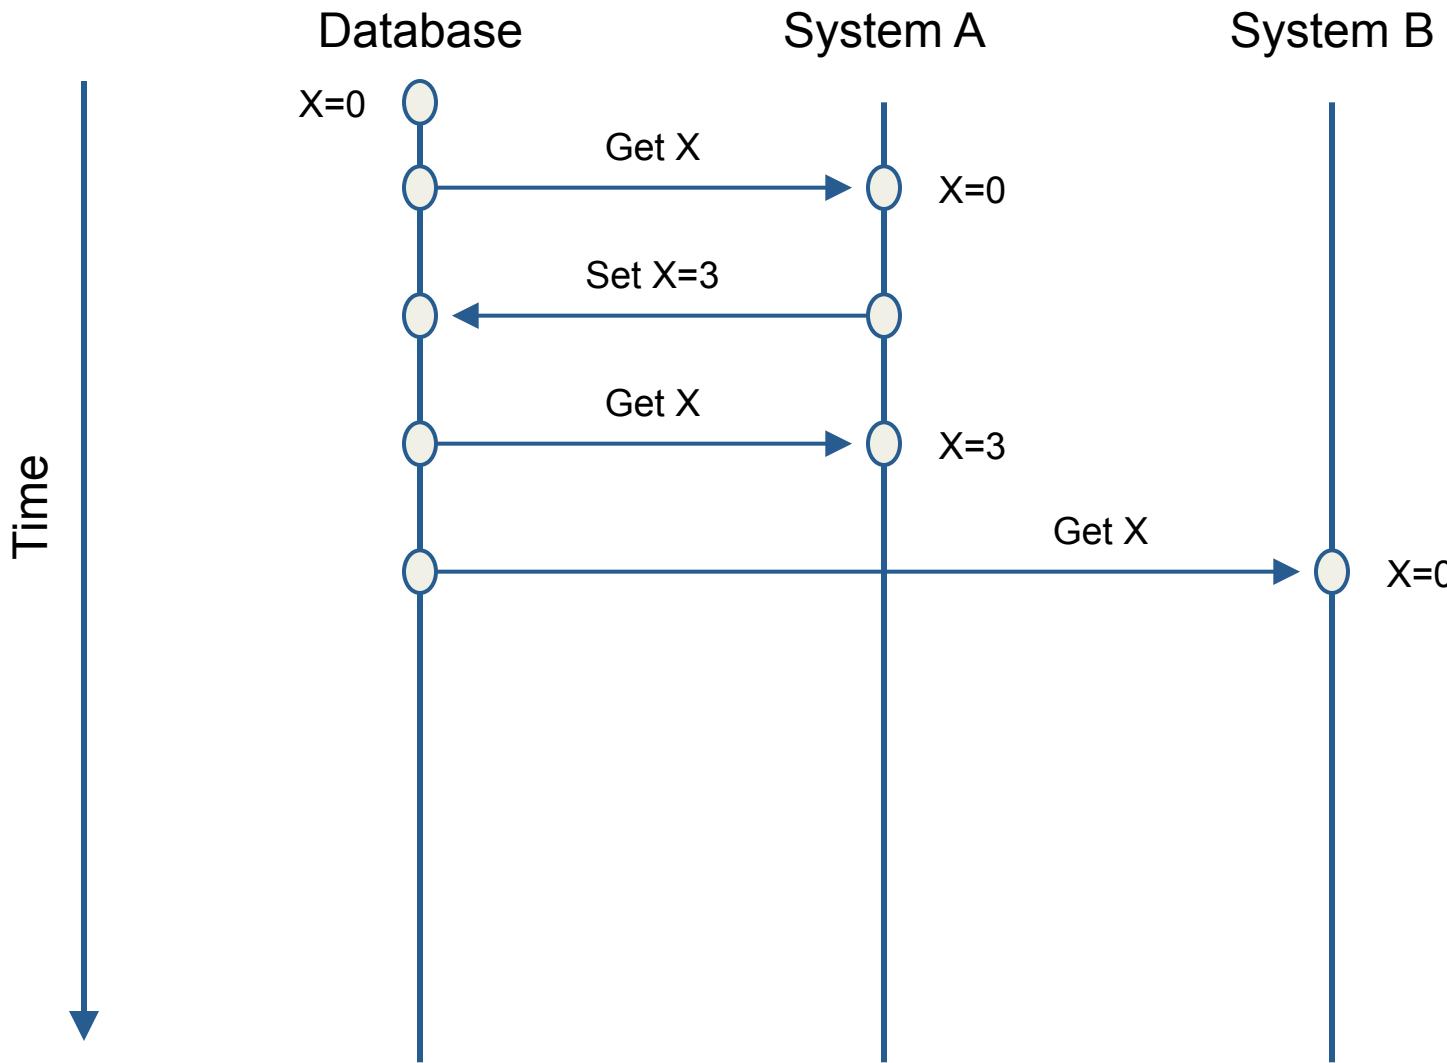
---



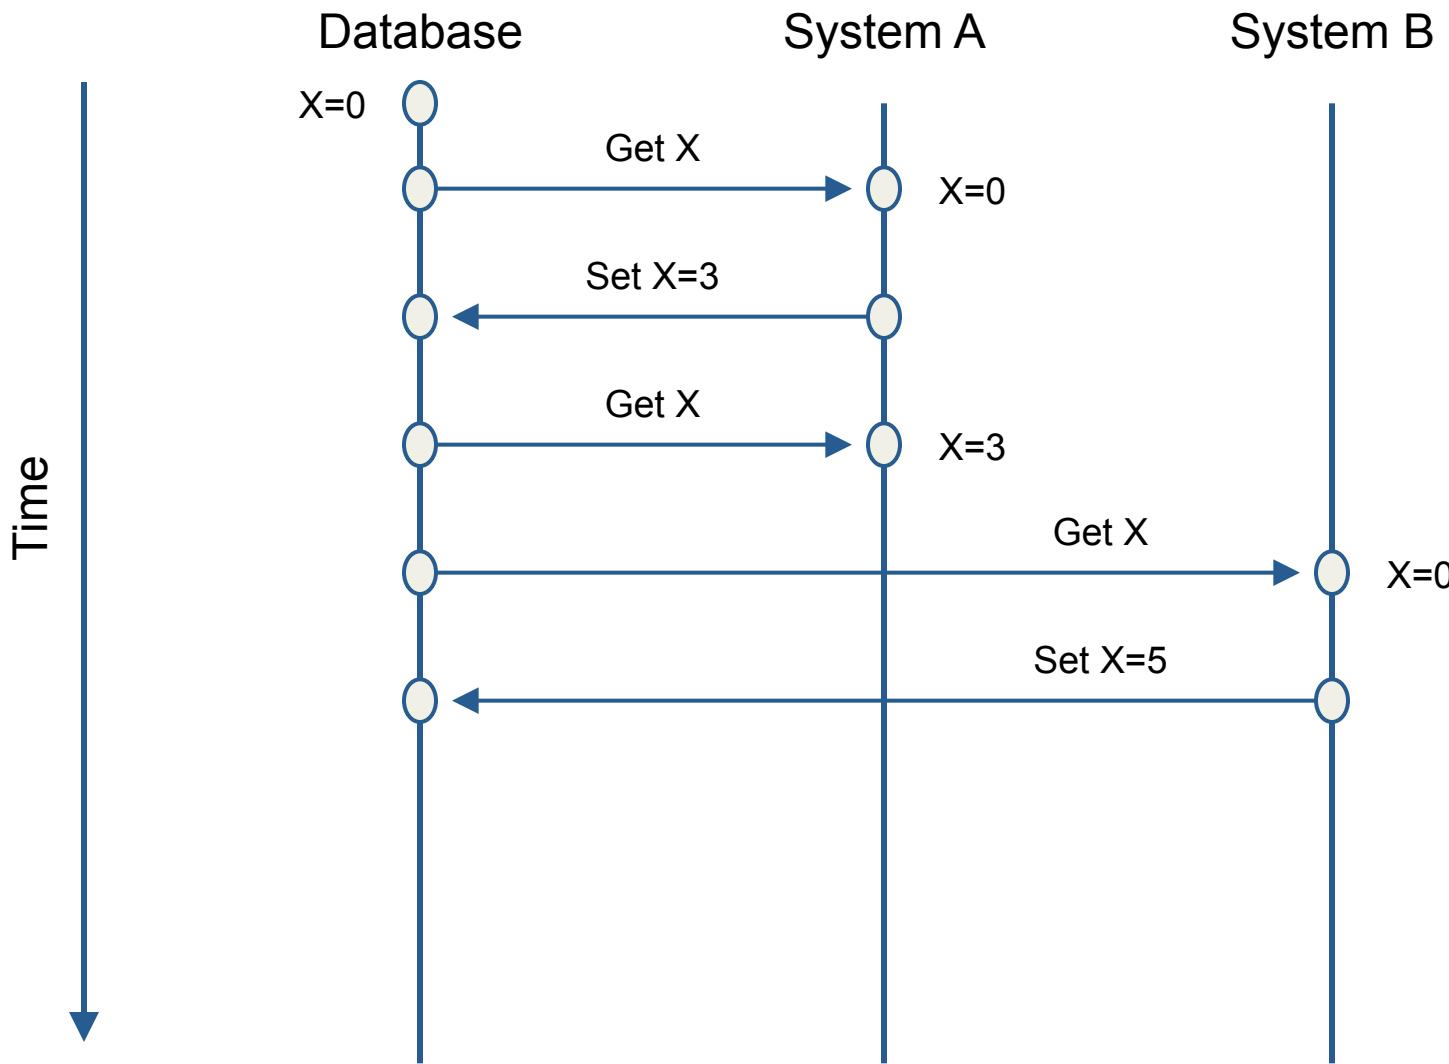
# Weak Consistency



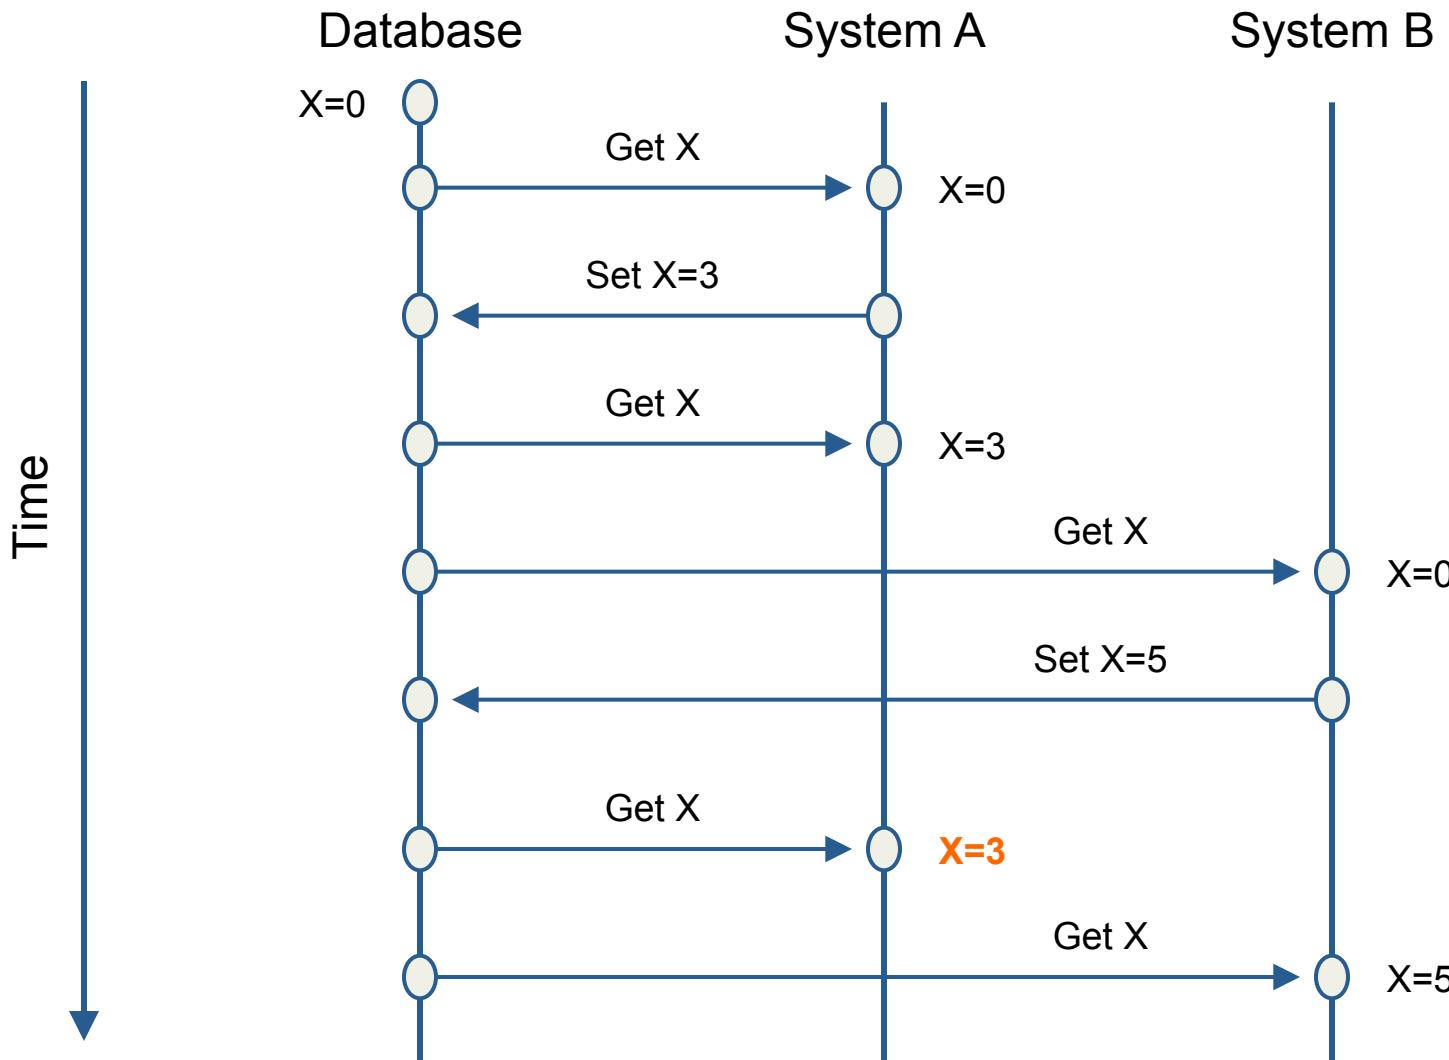
# Weak Consistency



# Weak Consistency



# Weak Consistency



# Consistency

---

- ◆ Does it matter if not consistent?
  - Results of search on Amazon....

# NoSQL Types

---

## Types

- ◆ Column Store
- ◆ Key-Value Store
- ◆ Document Store
- ◆ Graph Store

# Column Store

---

<Name> - <Value> - <Timestamp>

Andy	-	Boulder	-	140612
John	-	Boston	-	140612
Alice	-	Seattle	-	140613
Andy	-	Denver	-	140614
Stacy	-	Berlin	-	140615
John	-	New York	-	140615
Alice	-	Seattle	-	140616
Andy	-	Boulder	-	140615

# Column Store

---

J U U U U U U U U U U U U  
accumulo  
J N N N N N N N N N N N N

A P A C H E  
**HBASE**



# Key-Value Store

---

<Key>: <Value>

1675 : “Andy”

1676 : “John”

1677 : “Alice”

1678 : “Stacey”

1679 : “Dave”

1680 : “Rick”

1681 : “Jill”

# Key-Value Store

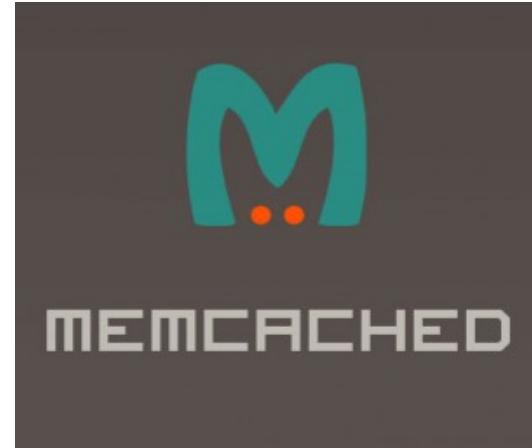
---



redis



FOUNDATION**DB**



# Document Store

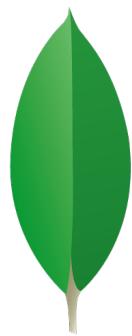
---

<Key>: <Document>

```
1675 : {Name: "Andy", Location: "Boulder", Activity: "School"}  
1676 : {Name: "John", Location: "Boston"}  
1677 : {Name: "Alice", Age: 27, Location: "Seattle"}  
1678 : {Name: "Stacey", Location: "Berlin", Activity: "Work"}  
1679 : {Name: "Dave", Age: 30, Activity: "Work"}  
1680 : {Name: "Rick", Location: "Denver"}  
1681 : {Name: "Jill", Activity: "Work"}
```

# Document Store

---



mongoDB



Cloudant

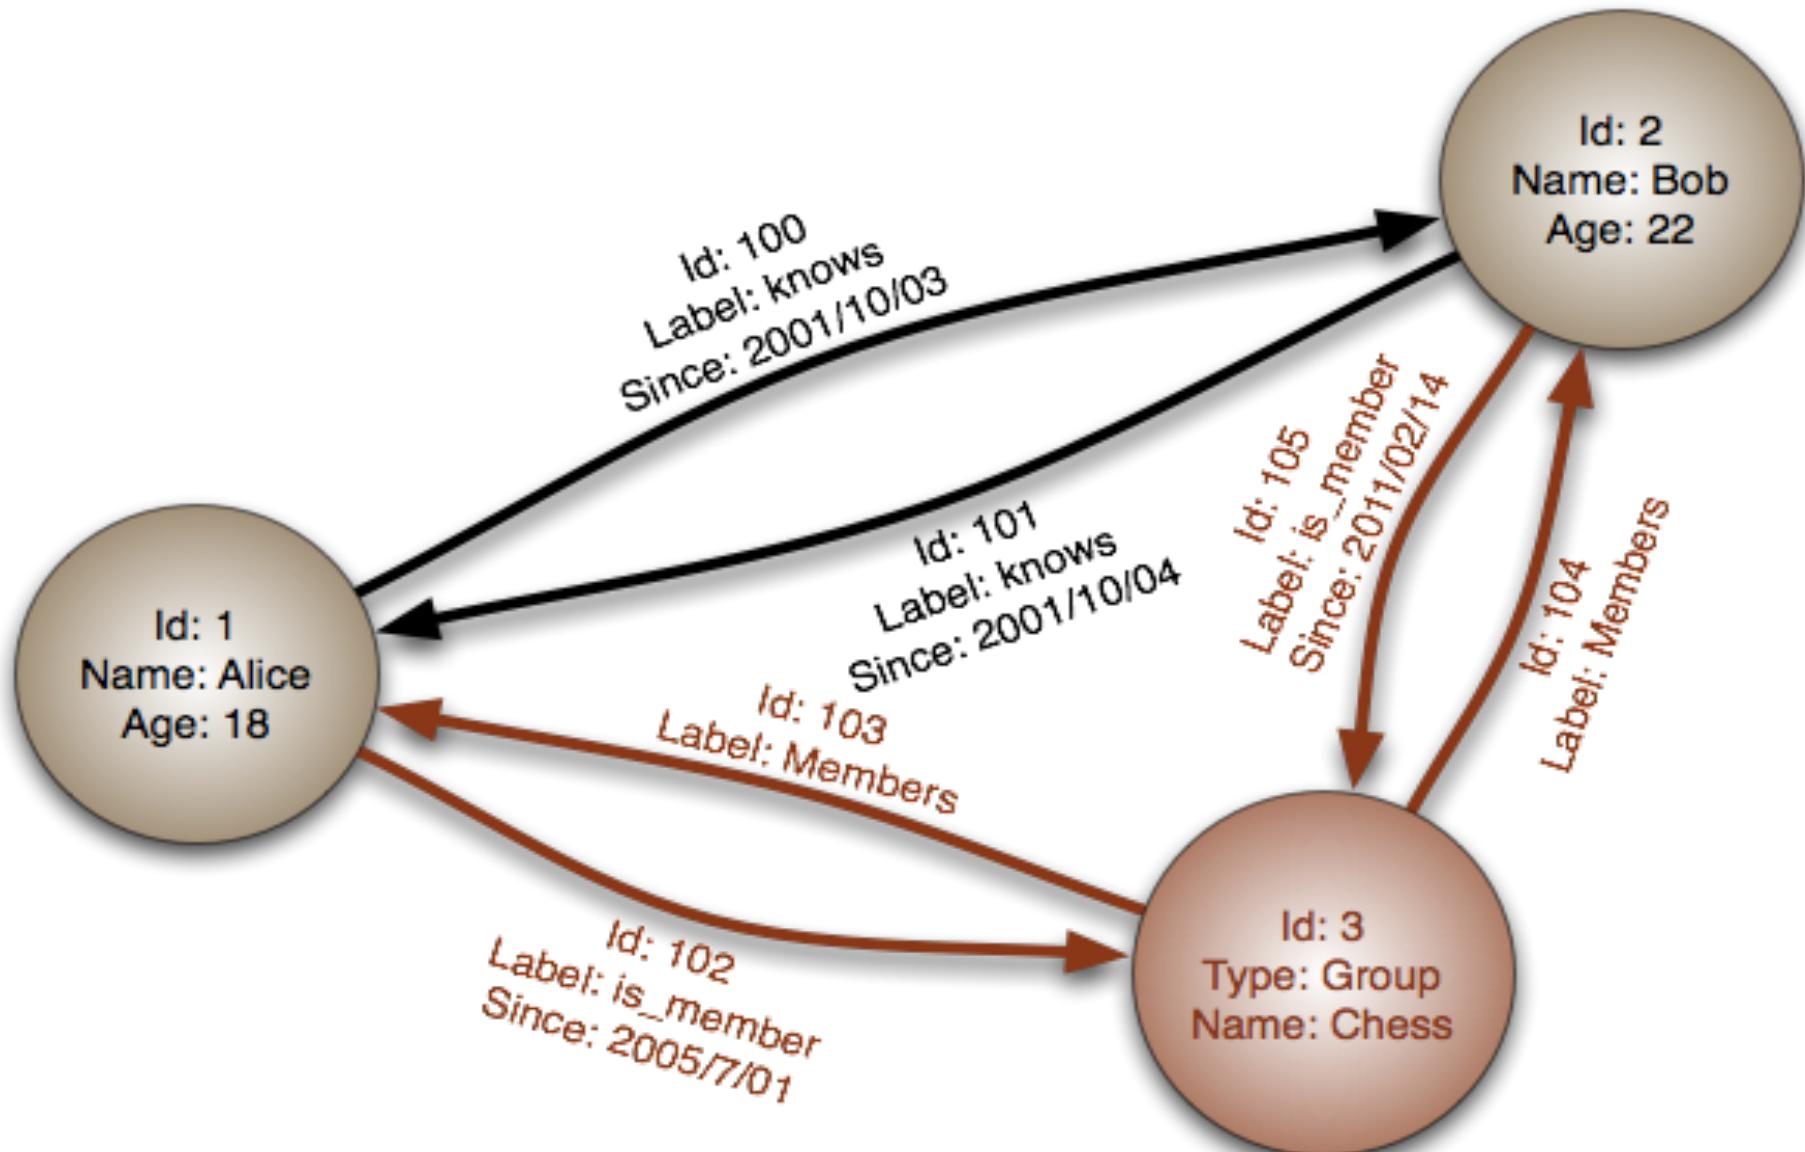


Couchbase



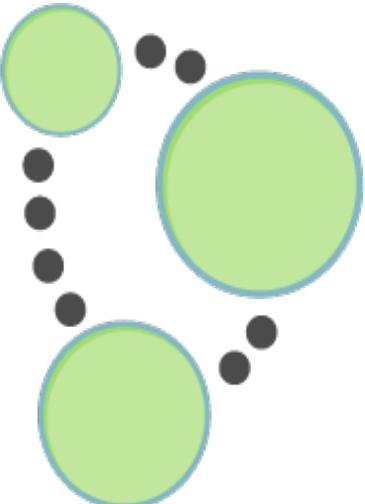
CouchDB  
relax

# Graph Store



# Graph Store

---



**Neo4j**



**ArangoDB**

# When to Use NoSQL

---

- ◆ Simpler Interface
- ◆ Loosely Structured Data
- ◆ Weaker Consistency
- ◆ Higher Performance
- ◆ Higher Availability
- ◆ Horizontal Scaling



- ◆ MongoDB (from "humongous")

- Open-source document database
- Written in C++
- Document-Oriented Storage

**BSON** { 01010100  
11101011  
10101110  
01010101 }

BSON [*bee · sahn*], short for Binary JSON, is a binary-encoded serialization of JSON-like documents. Like JSON, BSON supports the

BSON was designed to have three characteristics:

– [Document](#)

# mongoDB

---

## ◆ Examples SQL vs MongoDB

SQL SELECT Statements

```
SELECT *
```

```
FROM users
```

```
SELECT id,  
       user_id,  
       status  
  FROM users
```

MongoDB find() Statements

```
db.users.find()
```

```
db.users.find(  
  { },  
  { user_id: 1, status: 1 }  
)
```

## SQL SELECT Statements

## MongoDB find() Statements

```
SELECT *  
FROM users
```

```
db.users.find()
```

```
SELECT id,  
       user_id,  
       status  
FROM users
```

```
db.users.find(  
              { },  
              { user_id: 1, status: 1 }  
)
```

```
SELECT user_id, status  
FROM users
```

```
db.users.find(  
              { },  
              { user_id: 1, status: 1, _id: 0 }  
)
```

```
SELECT *  
FROM users  
WHERE status = "A"
```

```
db.users.find(  
              { status: "A" }  
)
```

```
SELECT user_id, status  
FROM users  
WHERE status = "A"
```

```
db.users.find(  
              { status: "A" },  
              { user_id: 1, status: 1, _id: 0 }  
)
```